

Introduction to Data Management

DATA 514

Unit 3: RDBMS Internals

Logical and Physical Plans
Query Execution
Query Performance & Optimization

Introduction to Data Management

DATA 514

Query Evaluation

Announcement

- **Midterm: next Tuesday, 5-6pm in class**
 - Open book (but you shouldn't need it)
 - Open computer: i.e., SQLite

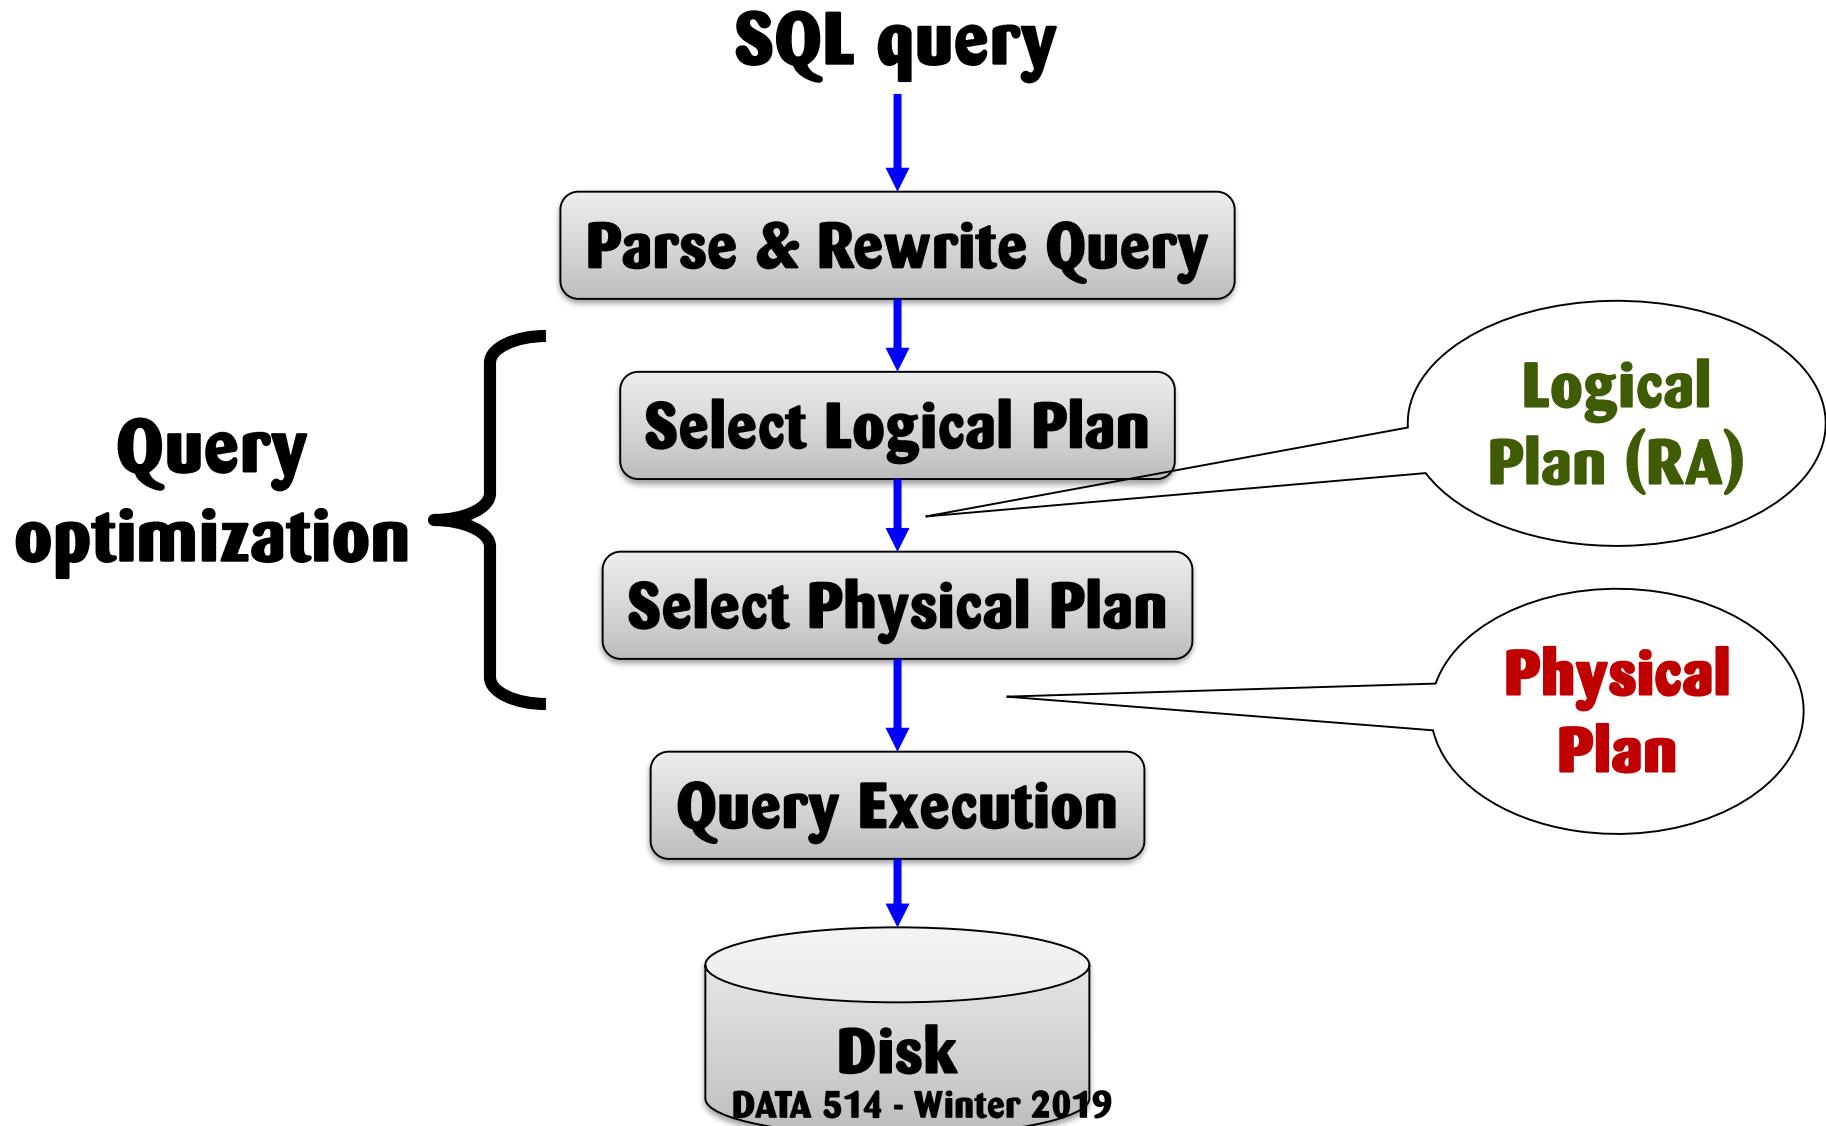
Class Overview

- **Unit 1: Intro**
- **Unit 2: Relational Data Models and Query Languages**
- **Unit 3: RDMBS internals and query optimization**
- **Unit 4: Parallel query processing**
- **Unit 5: DBMS usability, conceptual design**
- **Unit 6: Transactions**
- **Unit 7: Non-relational data**
- **Unit 8: Advanced topics (time permitting)**

From Logical RA Plans to Physical Plans

DATA 514 - Winter 2019

Query Evaluation Steps



Logical vs Physical Plans

- **Logical plans:**
 - Created by the parser from the input SQL text
 - Expressed as a Relational Algebra tree
 - Multiple possible logical plans
- **Physical plans:**
 - Goal is to choose an efficient implementation for each operation in the RA tree
 - Each logical plan has many possible physical plans

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Review: Relational Algebra

```
SELECT sname  
FROM Supplier x, Supply y  
WHERE x.sid = y.sid  
    and y.pno = 2  
    and x.scity = 'Seattle'  
    and x.sstate = 'WA'
```



Π_{sname}

$\sigma_{scity='Seattle' \text{ and } sstate='WA' \text{ and } pno=2}$



$sid = sid$

Supplier

Supply

Relational algebra expression is also called the “logical query plan”

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Physical Query Plan 1

A physical query plan is a logical query plan annotated with physical implementation details

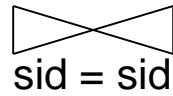
(On the fly)

π_{sname}

(On the fly)

$\sigma_{scity = \text{'Seattle'} \text{ and } sstate = \text{'WA'} \text{ and } pno = 2}$

(Nested loop)



Supplier
(Table scan)

Supply
(Table scan)

```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
    and y.pno = 2
    and x.scity = 'Seattle'
    and x.sstate = 'WA'
```

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Physical Query Plan 2

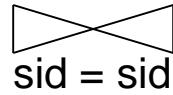
(On the fly)

π_{sname}

(On the fly)

$\sigma_{scity = \text{'Seattle'} \text{ and } sstate = \text{'WA'} \text{ and } pno = 2}$

(Hash join)



Supplier
(Table scan)

Supply
(Table scan)

**Same logical query plan
Different physical plan**

```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
and y.pno = 2
and x.scity = 'Seattle'
and x.sstate = 'WA'
```

`Supplier(sid, sname, scity, sstate)`

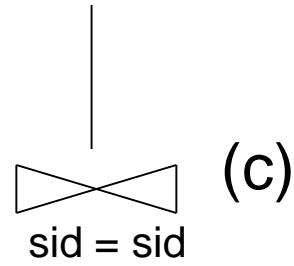
`Supply(sid, pno, quantity)`

Physical Query Plan 3

(On the fly)

π_{sname} (d)

(Sort-merge join)



(Scan & write to T1)

(a) $\sigma_{scity='Seattle' \text{ and } sstate='WA'}$

Supplier
(Table scan)

(b) $\sigma_{pno=2}$ (Scan & write to T2)

Supply
(Table scan)

Different but equivalent logical query plan; different physical plan

```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
and y.pno = 2
and x.scity = 'Seattle'
and x.sstate = 'WA'
```

Query Optimization Problem

- **For each SQL query... many logical plans**
- **For each logical plan... many physical plans**
- **Next: we will discuss physical operators;**

How exactly is a query executed?

Query Execution

Physical Operators

Each of the logical operators may have one or more implementations = physical operators

Will discuss several basic physical operators, with a focus on join

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Memory-resident Algorithms

- **Logical operation (JOIN):**

Supplier $\bowtie_{\text{sid}=\text{sid}}$ **Supply**

- **Assuming the tables are both resident in main memory, consider three physical options:**

Option	Execution time
Nested Loop JOIN	?
Merge JOIN	?
Hash JOIN	?

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Memory-resident Algorithms

- **Logical operation (JOIN):**

Supplier $\bowtie_{\text{sid}=\text{sid}}$ **Supply**

- **Assuming the tables are both resident in main memory, consider three physical options:**

Nested Loop JOIN

$O(n^2)$

Merge JOIN

$O(n \log n)$

Hash JOIN

$O(n \dots n^2)$

Hash Tables

Duplicates OK,
lead to **collisions**

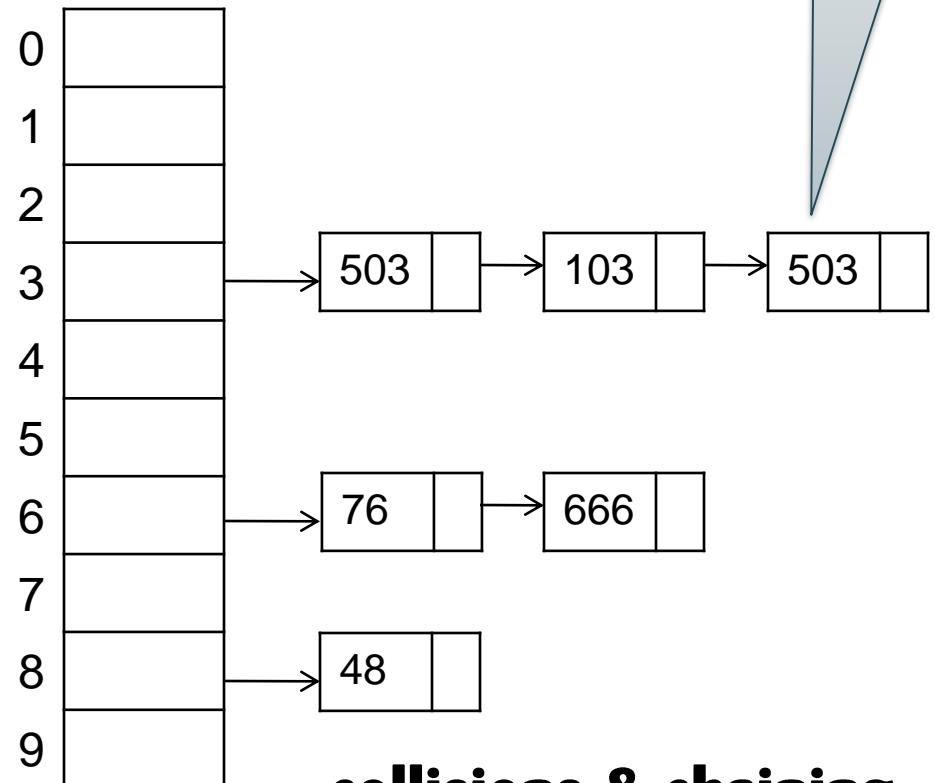
A (**naïve**) hash function:

$$h(x) = x \bmod 10$$

Operations:

$\text{find}(103) = ??$

$\text{insert}(488) = ??$



collisions & chaining

Hash Tables

- **insert(k , v) = inserts a key k with value v**
- **Many values for one key**
 - Hence, duplicate k 's are OK
- **find(k) = returns the list of all values, v associated with the key k**

Implementing Query Operators with the Iterator Interface

Each operator implements three methods:

- **open()**
- **next()**
- **close()**

Implementing Query Operators with the Iterator Interface

Example “on the fly” selection operator

```
interface Operator {
```

```
}
```

Implementing Query Operators with the Iterator Interface

Example “on the fly” selection operator

```
interface Operator {  
  
    // initializes operator state  
    // and sets parameters  
    void open (...);  
  
}  
}
```

Implementing Query Operators with the Iterator Interface

Example “on the fly” selection operator

```
interface Operator {  
  
    // initializes operator state  
    // and sets parameters  
    void open (...);  
  
    // calls next() on its inputs  
    // processes an input tuple  
    // produces output tuple(s)  
    // returns null when done  
    Tuple next ();  
}
```

Implementing Query Operators with the Iterator Interface

Example “on the fly” selection operator

```
interface Operator {  
  
    // initializes operator state  
    // and sets parameters  
    void open (...);  
  
    // calls next() on its inputs  
    // processes an input tuple  
    // produces output tuple(s)  
    // returns null when done  
    Tuple next ();  
  
    // cleans up (if any)  
    void close ();  
}
```

Implementing Query Operators with the Iterator Interface

Example “on the fly” selection operator

```
interface Operator {  
  
    // initializes operator state  
    // and sets parameters  
    void open (...);  
  
    // calls next() on its inputs  
    // processes an input tuple  
    // produces output tuple(s)  
    // returns null when done  
    Tuple next ();  
  
    // cleans up (if any)  
    void close ();  
}
```

```
class Select implements Operator {...  
    void open (Predicate p,  
              Operator child) {  
        this.p = p; this.child = child;  
    }  
}
```

Implementing Query Operators with the Iterator Interface

Example “on the fly” selection operator

```
interface Operator {  
  
    // initializes operator state  
    // and sets parameters  
    void open (...);  
  
    // calls next() on its inputs  
    // processes an input tuple  
    // produces output tuple(s)  
    // returns null when done  
    Tuple next ();  
  
    // cleans up (if any)  
    void close ();  
}
```

```
class Select implements Operator {...  
    void open (Predicate p,  
              Operator child) {  
        this.p = p; this.child = child;  
    }  
    Tuple next () {  
        ...  
    }  
}
```

Implementing Query Operators with the Iterator Interface

Example “on the fly” selection operator

```
interface Operator {  
  
    // initializes operator state  
    // and sets parameters  
    void open (...);  
  
    // calls next() on its inputs  
    // processes an input tuple  
    // produces output tuple(s)  
    // returns null when done  
    Tuple next ();  
  
    // cleans up (if any)  
    void close ();  
}
```

```
class Select implements Operator {...  
    void open (Predicate p,  
              Operator child) {  
        this.p = p; this.child = child;  
    }  
    Tuple next () {  
        boolean found = false;  
        Tuple r = null;  
        while (!found) {  
            r = child.next();  
            if (r == null) break;  
            found = p(in);  
        }  
    }  
}
```

Implementing Query Operators with the Iterator Interface

Example “on the fly” selection operator

```
interface Operator {  
  
    // initializes operator state  
    // and sets parameters  
    void open (...);  
  
    // calls next() on its inputs  
    // processes an input tuple  
    // produces output tuple(s)  
    // returns null when done  
    Tuple next ();  
  
    // cleans up (if any)  
    void close ();  
}
```

```
class Select implements Operator {...  
    void open (Predicate p,  
              Operator child) {  
        this.p = p; this.child = child;  
    }  
    Tuple next () {  
        boolean found = false;  
        Tuple r = null;  
        while (!found) {  
            r = child.next();  
            if (r == null) break;  
            found = p(in);  
        }  
        return r;  
    }  
}
```

Implementing Query Operators with the Iterator Interface

Example “on the fly” selection operator

```
interface Operator {  
  
    // initializes operator state  
    // and sets parameters  
    void open (...);  
  
    // calls next() on its inputs  
    // processes an input tuple  
    // produces output tuple(s)  
    // returns null when done  
    Tuple next ();  
  
    // cleans up (if any)  
    void close ();  
}
```

```
class Select implements Operator {...  
    void open (Predicate p,  
              Operator child) {  
        this.p = p; this.child = child;  
    }  
    Tuple next () {  
        boolean found = false;  
        Tuple r = null;  
        while (!found) {  
            r = child.next();  
            if (r == null) break;  
            found = p(in);  
        }  
        return r;  
    }  
    void close () { child.close(); }  
}
```

Implementing Query Operators with the Iterator Interface

```
interface Operator {  
  
    // initializes operator state  
    // and sets parameters  
    void open (...);  
  
    // calls next() on its inputs  
    // processes an input tuple  
    // produces output tuple(s)  
    // returns null when done  
    Tuple next ();  
  
    // cleans up (if any)  
    void close ();  
}
```

Query plan execution

```
Operator q = parse("SELECT ...");  
q = optimize(q);  
  
q.open();  
while (true) {  
    Tuple t = q.next();  
    if (t == null) break;  
    else printOnScreen(t);  
}  
q.close();
```

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Pipelining

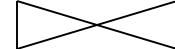
(On the fly)

Π_{sname}

(On the fly)

$\sigma_{\text{scity} = \text{'Seattle'} \text{ and } \text{sstate} = \text{'WA'} \text{ and } \text{pno} = 2}$

(Nested loop)


sno = sno

Suppliers
(Table scan)

Supplies
(Table scan)

Discuss:
open/next/close
for nested loop join

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Pipelining

(On the fly)

Π_{sname} **open()**

(On the fly)

$\sigma_{\text{scity} = \text{'Seattle'} \text{ and } \text{sstate} = \text{'WA'} \text{ and } \text{pno} = 2}$

(Nested loop)

\bowtie
sno = sno

Suppliers
(Table scan)

Supplies
(Table scan)

Discuss: open/next/close
for nested loop join

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

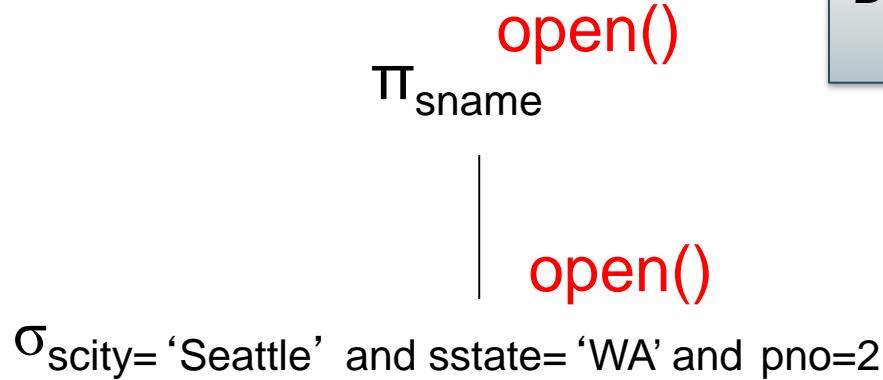
Pipelining

(On the fly)

(On the fly)

(Nested loop)

Suppliers
(Table scan)



Discuss: open/next/close
for nested loop join

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Pipelining

(On the fly)

Π_{sname} **open()**

(On the fly)

$\sigma_{\text{scity} = \text{'Seattle'} \text{ and } \text{sstate} = \text{'WA'} \text{ and } \text{pno} = 2}$

(Nested loop)

\bowtie
sno = sno **open()**

Suppliers
(Table **scan**)

Supplies
(Table **scan**)

Discuss: open/next/close
for nested loop join

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Pipelining

(On the fly)

(On the fly)

(Nested loop)

`open()`

Suppliers

(Table scan)

`open()`

`open()`

`open()`

\bowtie
 $sno = sno$

Discuss: open/next/close
for nested loop join

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Pipelining

(On the fly)

(On the fly)

(Nested loop)

`Suppliers
(Table scan)`

$\sigma_{scity = \text{'Seattle'} \text{ and } sstate = \text{'WA'} \text{ and } pno = 2}$

Π_{sname}

`open()`

`open()`

`open()`

\bowtie
 $sno = sno$

`Supplies
(Table scan)`

Discuss: open/next/close
for nested loop join

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Pipelining

(On the fly)

Π_{sname}
next()

(On the fly)

$\sigma_{scity = \text{'Seattle'} \text{ and } sstate = \text{'WA'} \text{ and } pno = 2}$

(Nested loop)

$sno = sno$

Suppliers
(Table scan)

Supplies
(Table scan)

Discuss: open/next/close
for nested loop join

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Pipelining

(On the fly)

Π_{sname}

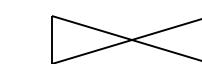
next()

(On the fly)

$\sigma_{scity='Seattle' \text{ and } sstate='WA' \text{ and } pno=2}$

next()

(Nested loop)


sno = sno

Suppliers
(Table scan)

Supplies
(Table scan)

Discuss: open/next/close
for nested loop join

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Pipelining

(On the fly)

Π_{sname}

next()

(On the fly)

$\sigma_{scity='Seattle' \text{ and } sstate='WA' \text{ and } pno=2}$

next()

(Nested loop)

 sno = sno

next()

Suppliers
(Table scan)

Supplies
(Table scan)

Discuss: open/next/close
for nested loop join

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Pipelining

(On the fly)

Π_{sname}

next()

(On the fly)

$\sigma_{scity='Seattle' \text{ and } sstate='WA' \text{ and } pno=2}$

next()

(Nested loop)

\bowtie
sno = sno

next()

next()

Suppliers

(Table scan)

Supplies

(Table scan)

Discuss: open/next/close
for nested loop join

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Pipelining

(On the fly)

Π_{sname}

next()

(On the fly)

$\sigma_{scity='Seattle' \text{ and } sstate='WA' \text{ and } pno=2}$

next()

(Nested loop)

\bowtie
sno = sno

next()

next()

Suppliers

(Table scan)

Discuss: open/next/close
for nested loop join

next()

Supplies

(Table scan)

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Pipelining

(On the fly)

Π_{sname}

next()

(On the fly)

$\sigma_{scity='Seattle' \text{ and } sstate='WA' \text{ and } pno=2}$

next()

(Nested loop)

\bowtie
sno = sno

next()

next()

Suppliers

(Table scan)

next()
next()

Supplies
(Table scan)

Discuss: open/next/close
for nested loop join

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Pipelining

(On the fly)

Π_{sname}

Discuss hash-join
in class

(On the fly)

$\sigma_{scity = \text{'Seattle'} \text{ and } sstate = \text{'WA'} \text{ and } pno = 2}$

(Hash Join)

 $sno = sno$

Suppliers
(Table scan)

Supplies
(Table scan)

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Pipelining

(On the fly)

Π_{sname}

Discuss hash-join
in class

(On the fly)

$\sigma_{scity = \text{'Seattle'} \text{ and } sstate = \text{'WA'} \text{ and } pno = 2}$

(Hash Join)

Tuples from
here are
pipelined


 $sno = sno$

Suppliers
(Table scan)

Supplies
(Table scan)

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Pipelining

(On the fly)

Π_{sname}

Discuss hash-join
in class

(On the fly)

$\sigma_{scity = \text{'Seattle'} \text{ and } sstate = \text{'WA'} \text{ and } pno = 2}$

(Hash Join)

Tuples from here are pipelined

Suppliers
(Table scan)

$sno = sno$

Tuples from here are “blocked”

Supplies
(Table scan)

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Blocked Execution

(On the fly)

Π_{sname}

(On the fly)

$\sigma_{scity = \text{'Seattle'} \text{ and } sstate = \text{'WA'} \text{ and } pno = 2}$

(Merge Join)


 $sno = sno$

Suppliers
(Table scan)

Supplies
(Table scan)

Discuss merge-join
in class

`Supplier(sid, sname, scity, sstate)`

`Supply(sid, pno, quantity)`

Blocked Execution

(On the fly)

Π_{sname}

Discuss merge-join
in class

(On the fly)

$\sigma_{scity='Seattle' \text{ and } sstate='WA' \text{ and } pno=2}$

(Merge Join)

$sno = sno$

Blocked

Blocked

Suppliers
(Table scan)

Supplies
(Table scan)

Pipelined Execution

- **Tuples generated by an operator are immediately sent to the parent**
- **Benefits:**
 - No operator synchronization issues
 - No need to buffer tuples between operators
 - Saves cost of writing intermediate data to disk
 - Saves cost of reading intermediate data from disk
- **This approach is used whenever possible**

Query Execution Bottom Line

- SQL query transformed into physical plan
 - Access path selection for each relation
 - Scan the relation or use an index (next lecture)
 - Implementation choice for each operator
 - Nested loop join, hash join, etc.
 - Scheduling decisions for operators
 - Pipelined execution or intermediate materialization
- Pipelined execution of physical plan

Recall: Physical Data Independence

- Applications are insulated from changes in physical storage details
- SQL and relational algebra facilitate physical data independence
 - Both languages input and output relations
 - Can choose different implementations for operators

Introduction to Database Systems

DATA 514

Lecture 15-16:
Basics of Data Storage and Indexes

Announcements

- **HW4 (datalog) due tomorrow (Tuesday)**
- **Midterm: Wednesday**
- **No sections on Thursday!**
- **HW5 (SQL++) due next Tuesday**

Query Performance

- **My database application is too slow... why?**
- **One of the queries is very slow... why?**
- **To understand performance, we need to understand:**
 - **How is data organized on disk**
 - **How to estimate query costs**
 - **In this course the focus is on disk-based DBMSs**

Data Storage

- DBMSs store data in files
 - Most common organization is row-wise storage
- On disk, the file(s) containing the database is organized into fixed-size blocks
 - Each block contains a set of tuples
 - example shows 4 blocks with 2 tuples each

ID	fName	IName
10	Tom	Hanks
20	Amy	Hanks
...		

10	Tom	Hanks	block 1
20	Amy	Hanks	
50	block 2
200	...		
220			block 3
240			
420			
800			

Data File Types

The data file can be one of:

- **Heap file**
 - **Unsorted**
- **Sequential file**
 - **Sorted according to some attribute(s) called key**

ID	fName	IName
10	Tom	Hanks
20	Amy	Hanks
...		

Data File Types

The data file can be one of:

- **Heap file**
 - Unsorted
- **Sequential file**
 - Sorted according to some attribute(s) called **key**

ID	fName	IName
10	Tom	Hanks
20	Amy	Hanks
...		

Note: key here means something different from primary key: it just means that we order the file according to that attribute. In our example we ordered by **ID**. Might as well order by **fName**, if that seems a better idea for the applications running on our database.

Index

- An additional file, that allows fast access to records in the data file given a search key

Index

- An additional file, that allows fast access to records in the data file given a search key
- The index contains (key, value) pairs:
 - The key = an attribute value (e.g., student ID or name)
 - The value = a pointer to the record

Index

- An additional file, that allows fast access to records in the data file given a search key
- The index contains (key, value) pairs:
 - The key = an attribute value (e.g., student ID or name)
 - The value = a pointer to the record
- Could have many indexes for one table

Here, Key = search key

This Is Not A Key



Different keys:

- Primary key – uniquely identifies a tuple
- Key of the sequential file – how the data file is sorted, if at all
- Index key – how the index is organized



This is not a pipe.

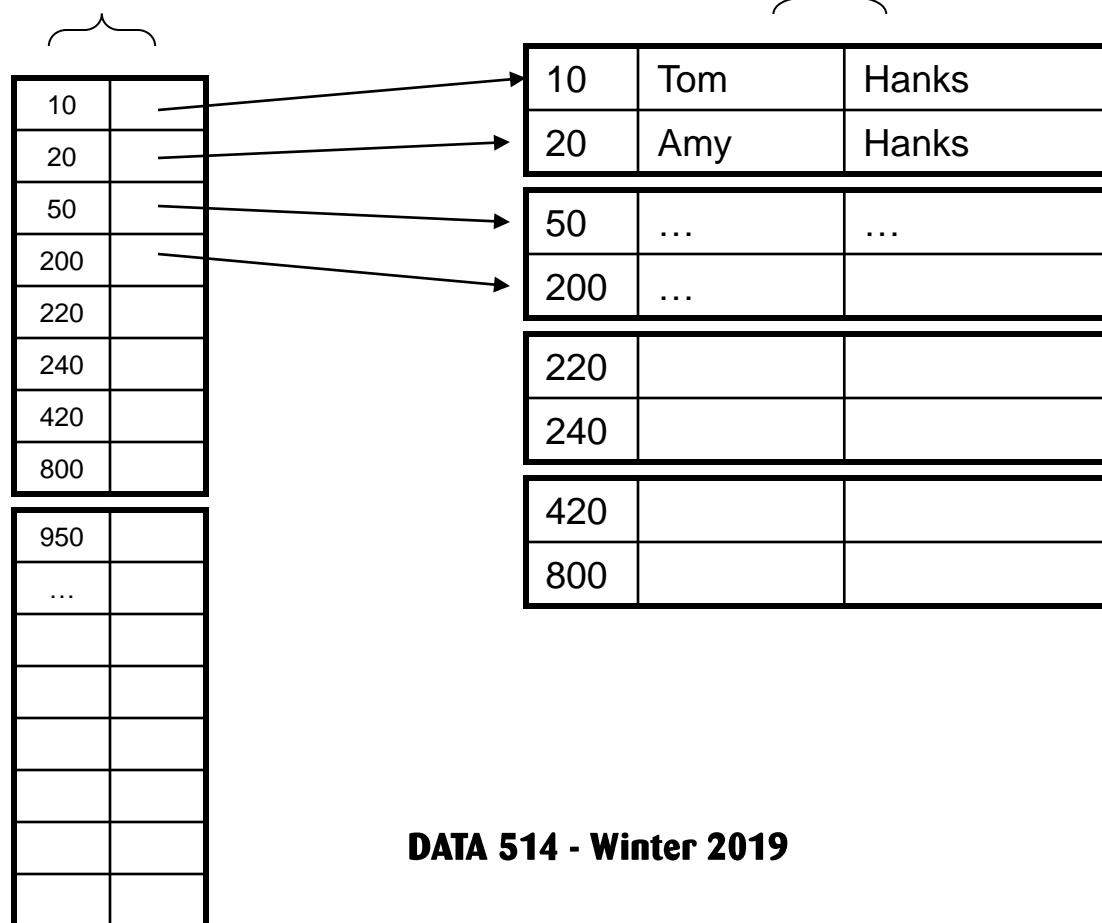
DATA 514 - Winter 2019



Example 1: Index on ID

Index **Student_ID** on **Student.ID**

Data File **Student**

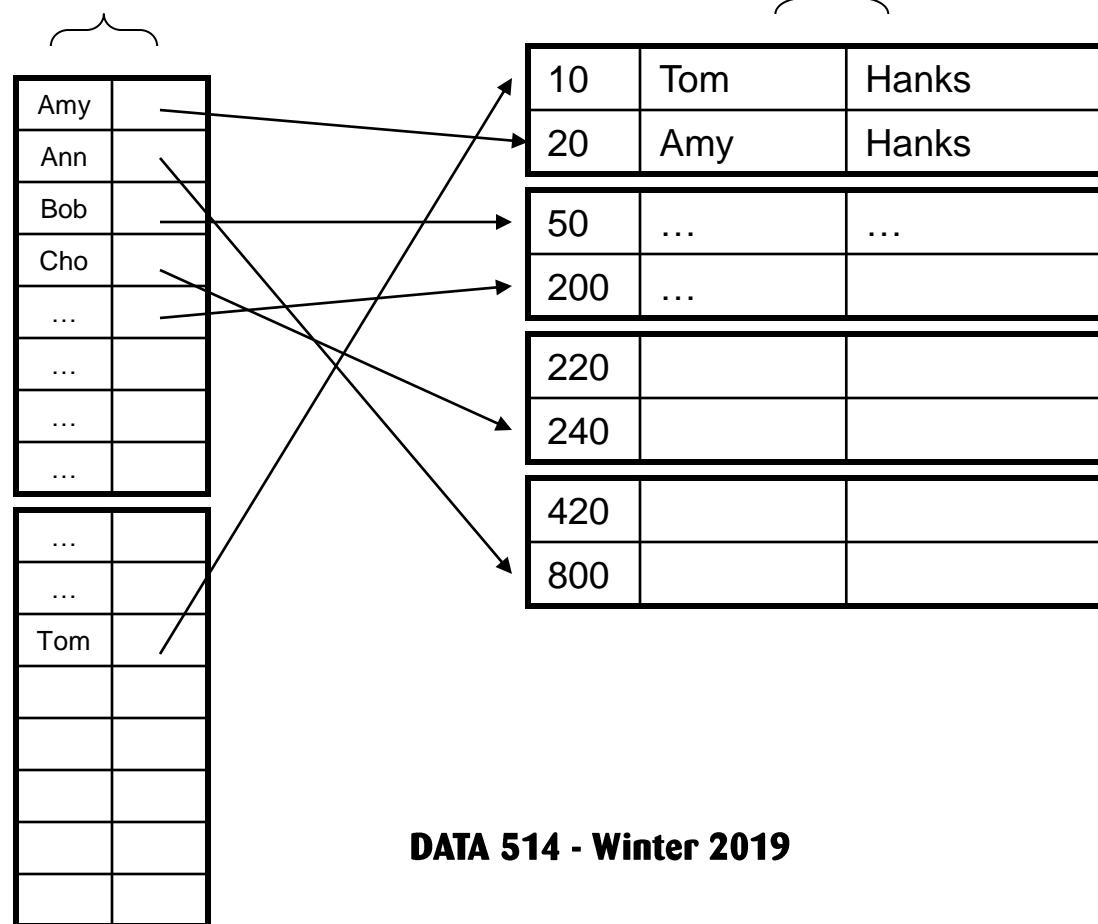


Example 2:

Index on fName

Index **Student_fName**
on **Student.fName**

Data File **Student**



Index Organization

We need a way to represent indexes after loading into memory so that they can be used

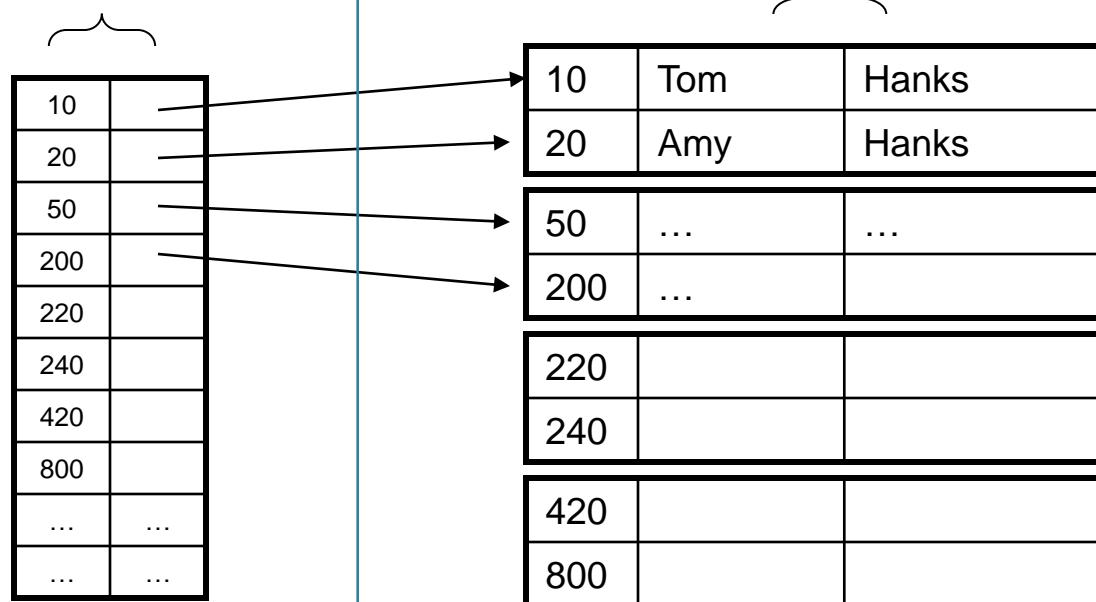
Several ways to do this:

- Hash table
- B+ trees – most popular
 - They are search trees, but they are not binary instead have higher fan-out
 - Will discuss them briefly next
- Specialized indexes: bit maps, R-trees, inverted index

ID	fName	IName
10	Tom	Hanks
20	Amy	Hanks
...		

Hash table example

Index **Student_ID** on **Student.ID**



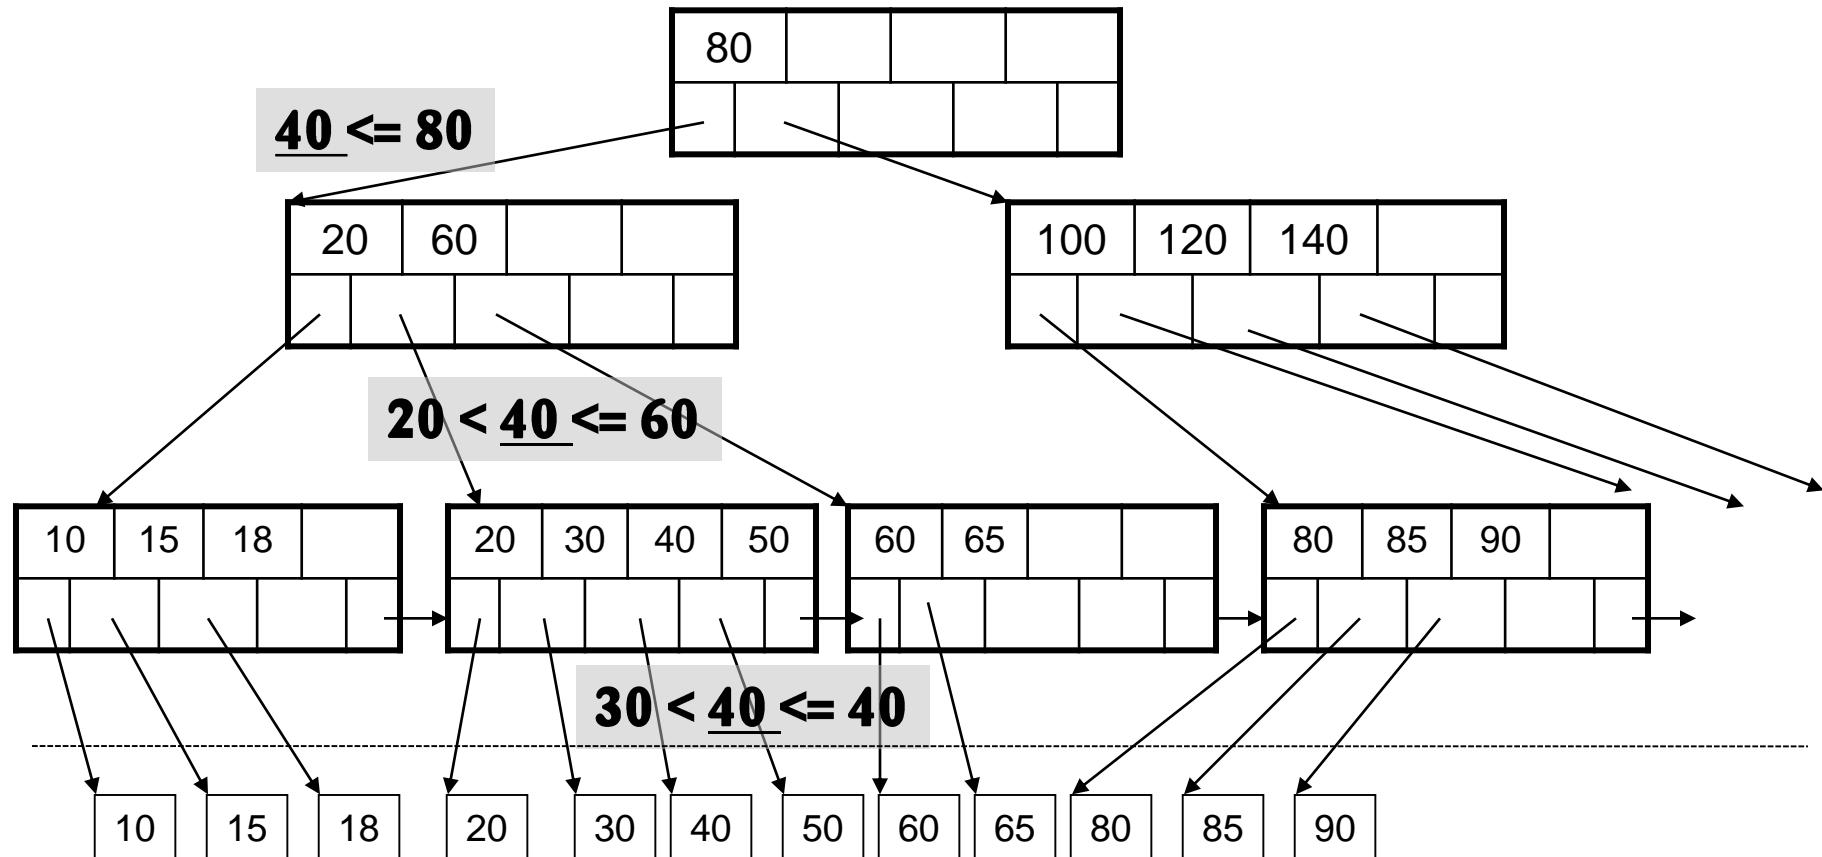
Index File
(preferably
in memory)

Data file
(on disk)

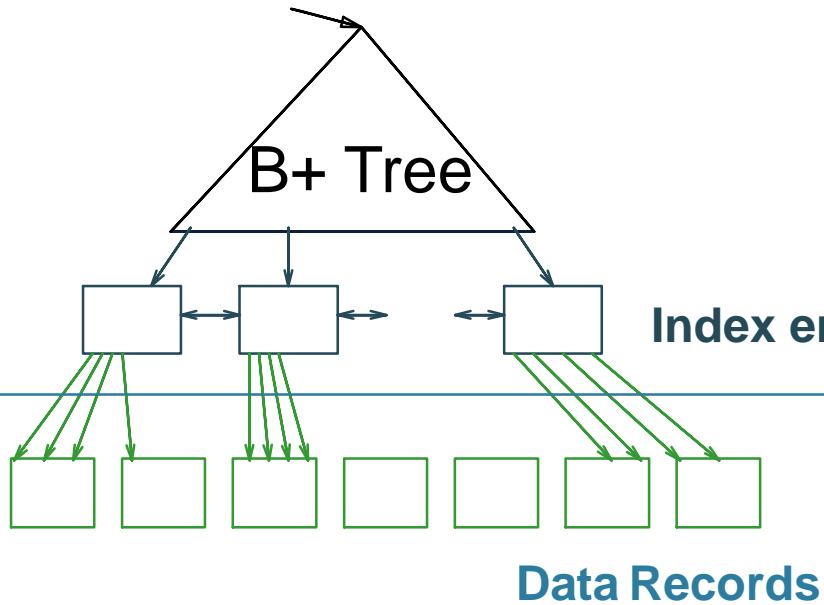
B+ Tree Index

Depth = 2

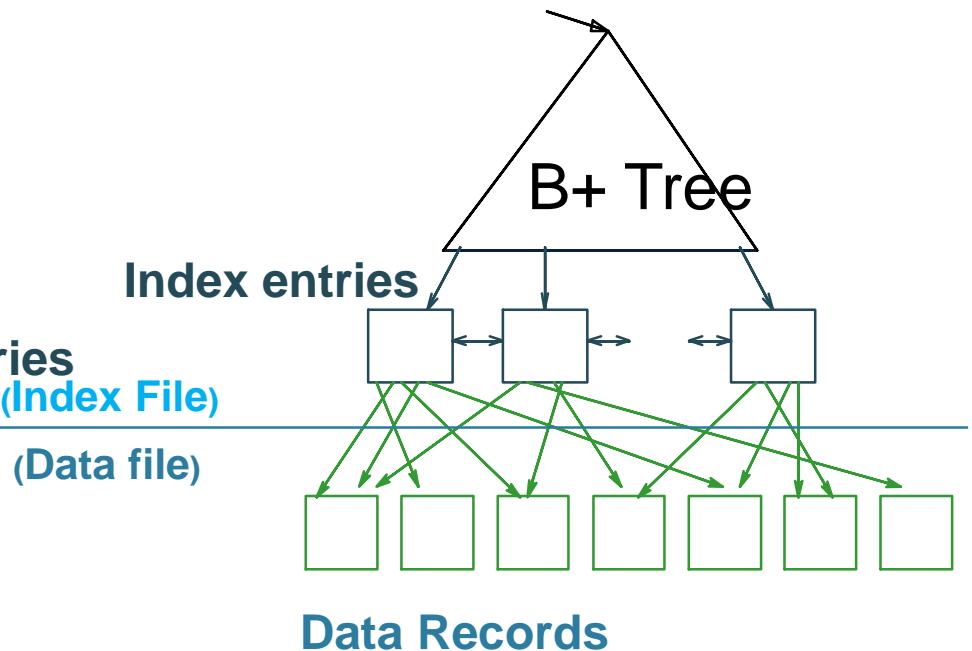
Find the key 40



Clustered vs Unclustered



CLUSTERED



UNCLUSTERED

Every table can have only one clustered index, but may have many unclustered indexes. Why?

Index Classification

- **Clustered/unclustered**
 - **Clustered = records close in index are close in data**
 - Option 1: Data inside data file is sorted on disk
 - Option 2: Store data directly inside the index (no separate files)
 - **Unclustered = records close in index may be far in data**

Index Classification

- **Clustered/unclustered**
 - Clustered = records close in index are close in data
 - Option 1: Data inside data file is sorted on disk
 - Option 2: Store data directly inside the index (no separate files)
 - Unclustered = records close in index may be far in data
- **Primary/secondary**
 - Meaning 1:
 - Primary = is over attributes that include the primary key
 - Secondary = otherwise
 - Meaning 2: means the same as clustered/unclustered

Index Classification

- **Clustered/unclustered**
 - Clustered = records close in index are close in data
 - Option 1: Data inside data file is sorted on disk
 - Option 2: Store data directly inside the index (no separate files)
 - Unclustered = records close in index may be far in data
- **Primary/secondary**
 - Meaning 1:
 - Primary = is over attributes that include the primary key
 - Secondary = otherwise
 - Meaning 2: means the same as clustered/unclustered
- **Organization: B+ tree or Hash table**

Scanning a Data File

- **Disks are mechanical devices!**
 - Technology from the 60s; density much higher now
- **Read only at the rotation speed!**
- **Consequence:**

Sequential scan is MUCH FASTER than random reads

 - **Good:** read blocks 1,2,3,4,5,...
 - **Bad:** read blocks 2342, 11, 321,9, ...
- **Rule of thumb:**
 - Random reading 1-2% of the file \approx sequential scanning the entire file; this is decreasing over time (because of increased density of disks)
- **Solid state (SSD): \$\$\$ expensive; put indexes, other “hot” data there, still too expensive for everything**



Summary So Far

- **Index = a file that enables direct access to records in another data file**
 - **B+ tree / Hash table**
 - **Clustered/unclustered**
- **Data resides on disk**
 - **Organized in blocks**
 - **Sequential reads are efficient**
 - **Random access less efficient**
 - **Random read 1-2% of data worse than sequential**

Student(ID, fname, lname)

Takes(studentID, courseID)

SELECT *

FROM Student x, Takes y

WHERE x.ID=y.studentID AND y.courseID > 300

Example

Student(ID, fname, lname)
Takes(studentID, courseID)

```
SELECT *
FROM Student x, Takes y
WHERE x.ID=y.studentID AND y.courseID > 300
```

Example

```
for y in Takes
  if courseID > 300 then
    for x in Student
      if x.ID=y.studentID
        output *
```

Student(ID, fname, lname)
Takes(studentID, courseID)

```
SELECT *
FROM Student x, Takes y
WHERE x.ID=y.studentID AND y.courseID > 300
```

Example

```
for y in Takes
  if courseID > 300 then
    for x in Student
      if x.ID=y.studentID
        output *
```

Assume the database has indexes on these attributes:

- **Takes_courseID** = index on Takes.courseID
- **Student_ID** = index on Student.ID

```
Student(ID, fname, lname)  
Takes(studentID, courseID)
```

```
SELECT *  
FROM Student x, Takes y  
WHERE x.ID=y.studentID AND y.courseID > 300
```

Example

```
for y in Takes  
if courseID > 300 then  
    for x in Student  
        if x.ID=y.studentID  
            output *
```

Assume the database has indexes on these attributes:

- **Takes_courseID** = index on Takes.courseID
- **Student_ID** = index on Student.ID

```
for y' in Takes_courseID where y'.courseID > 300
```

Student(ID, fname, lname)
Takes(studentID, courseID)

```
SELECT *
FROM Student x, Takes y
WHERE x.ID=y.studentID AND y.courseID > 300
```

Example

```
for y in Takes
  if courseID > 300 then
    for x in Student
      if x.ID=y.studentID
        output *
```

Assume the database has indexes on these attributes:

- **Takes_courseID** = index on Takes.courseID
- **Student_ID** = index on Student.ID

Index selection

```
for y' in Takes_courseID where y'.courseID > 300
  y = fetch the Takes record pointed to by y'
```

```
Student(ID, fname, lname)  
Takes(studentID, courseID)
```

```
SELECT *  
FROM Student x, Takes y  
WHERE x.ID=y.studentID AND y.courseID > 300
```

Example

```
for y in Takes  
if courseID > 300 then  
  for x in Student  
    if x.ID=y.studentID  
      output *
```

Assume the database has indexes on these attributes:

- **Takes_courseID** = index on Takes.courseID
- **Student_ID** = index on Student.ID

Index selection

Index join

```
for y' in Takes_courseID where y'.courseID > 300  
y = fetch the Takes record pointed to by y'  
for x' in Student_ID where x'.ID = y.studentID  
x = fetch the Student record pointed to by x'
```

```
Student(ID, fname, lname)  
Takes(studentID, courseID)
```

```
SELECT *  
FROM Student x, Takes y  
WHERE x.ID=y.studentID AND y.courseID > 300
```

Example

```
for y in Takes  
  if courseID > 300 then  
    for x in Student  
      if x.ID=y.studentID  
        output *
```

Assume the database has indexes on these attributes:

- **Takes_courseID** = index on Takes.courseID
- **Student_ID** = index on Student.ID

Index selection

Index join

```
for y' in Takes_courseID where y'.courseID > 300  
  y = fetch the Takes record pointed to by y'  
  for x' in Student_ID where x'.ID = y.studentID  
    x = fetch the Student record pointed to by x'  
    output *
```

Student(ID, fname, lname)
Takes(studentID, courseID)

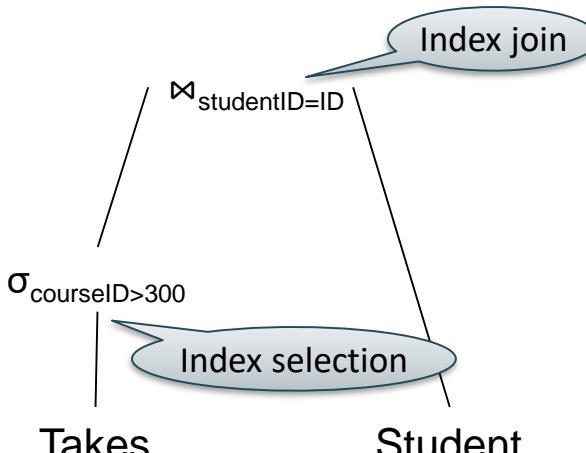
```
SELECT *
FROM Student x, Takes y
WHERE x.ID=y.studentID AND y.courseID > 300
```

Example

```
for y in Takes
  if courseID > 300 then
    for x in Student
      if x.ID=y.studentID
        output *
```

Assume the database has indexes on these attributes:

- **Takes_courseID** = index on Takes.courseID
- **Student_ID** = index on Student.ID



Index selection

```
for y' in Takes_courseID where y'.courseID > 300
  y = fetch the Takes record pointed to by y'
  for x' in Student_ID where x'.ID = y.studentID
    x = fetch the Student record pointed to by x'
    output *
```

Getting Practical: Creating Indexes in SQL

```
CREATE TABLE V(M int, N varchar(20), P int);
```

```
CREATE INDEX V1 ON V(N)
```

```
CREATE INDEX V2 ON V(P, M)
```

```
CREATE INDEX V3 ON V(M, N)
```

```
CREATE UNIQUE INDEX V4 ON V(N)
```

```
CREATE CLUSTERED INDEX V5 ON V(N)
```

Getting Practical: Creating Indexes in SQL

```
CREATE TABLE V(M int, N varchar(20), P int);
```

```
CREATE INDEX V1 ON V(N)
```

```
CREATE INDEX V2 ON V(P, M)
```

What does this mean?

```
CREATE INDEX V3 ON V(M, N)
```

```
CREATE UNIQUE INDEX V4 ON V(N)
```

```
CREATE CLUSTERED INDEX V5 ON V(N)
```

Getting Practical: Creating Indexes in SQL

```
CREATE TABLE V(M int, N varchar(20), P int);
```

yes

```
CREATE INDEX V1 ON V(N)
```

```
select *  
from V  
where P=55 and M=77
```

```
CREATE INDEX V2 ON V(P, M)
```

What does this mean?

```
CREATE INDEX V3 ON V(M, N)
```

```
CREATE UNIQUE INDEX V4 ON V(N)
```

```
CREATE CLUSTERED INDEX V5 ON V(N)
```

Getting Practical: Creating Indexes in SQL

```
CREATE TABLE V(M int, N varchar(20), P int);
```

yes

```
CREATE INDEX V1 ON V(N)
```

```
select *  
from V  
where P=55 and M=77
```

```
CREATE INDEX V2 ON V(P, M)
```

What does this mean?

```
CREATE INDEX V3 ON V(M, N)
```

```
select *  
from V  
where P=55
```

```
CREATE UNIQUE INDEX V4 ON V(N)
```

```
CREATE CLUSTERED INDEX V5 ON V(N)
```

Getting Practical: Creating Indexes in SQL

```
CREATE TABLE V(M int, N varchar(20), P int);
```

yes

```
CREATE INDEX V1 ON V(N)
```

```
select *  
from V  
where P=55 and M=77
```

```
CREATE INDEX V2 ON V(P, M)
```

What does this mean?

```
CREATE INDEX V3 ON V(M, N)
```

```
select *  
from V  
where P=55
```

```
CREATE UNIQUE INDEX V4 ON V(N)
```

yes

```
CREATE CLUSTERED INDEX V5 ON V(N)
```

Getting Practical: Creating Indexes in SQL

```
CREATE TABLE V(M int, N varchar(20), P int);
```

yes

```
CREATE INDEX V1 ON V(N)
```

```
select *  
from V  
where P=55 and M=77
```

```
CREATE INDEX V2 ON V(P, M)
```

What does this mean?

```
CREATE INDEX V3 ON V(M, N)
```

```
select *  
from V  
where P=55
```

```
CREATE UNIQUE INDEX V4 ON V(N)
```

```
select *  
from V  
where M=77
```

```
CREATE CLUSTERED INDEX V5 ON V(N)
```

Getting Practical: Creating Indexes in SQL

```
CREATE TABLE V(M int, N varchar(20), P int);
```

yes

```
CREATE INDEX V1 ON V(N)
```

```
select *  
from V  
where P=55 and M=77
```

```
CREATE INDEX V2 ON V(P, M)
```

What does this mean?

```
CREATE INDEX V3 ON V(M, N)
```

```
select *  
from V  
where P=55
```

```
CREATE UNIQUE INDEX V4 ON V(N)
```

```
select *  
from V  
where M=77
```

```
CREATE CLUSTERED INDEX V5 ON V(N)
```

yes

no

Getting Practical: Creating Indexes in SQL

```
CREATE TABLE V(M int, N varchar(20), P int);
```

yes

```
CREATE INDEX V1 ON V(N)
```

```
select *  
from V  
where P=55 and M=77
```

```
CREATE INDEX V2 ON V(P, M)
```

What does this mean?

```
CREATE INDEX V3 ON V(M, N)
```

```
select *  
from V  
where P=55
```

```
CREATE UNIQUE INDEX V4 ON V(N)
```

```
select *  
from V  
where M=77
```

```
CREATE CLUSTERED INDEX V5 ON V(N)
```

Not supported
in SQLite

Which Indexes?

ID	fName	IName
10	Tom	Hanks
20	Amy	Hanks
...		

- How many indexes could we create?
- Which indexes should we create?

Which Indexes?

ID	fName	IName
10	Tom	Hanks
20	Amy	Hanks
...		

- How many indexes could we create?
- Which indexes should we create?

This is a very hard problem!

Which Indexes?

ID	fName	IName
10	Tom	Hanks
20	Amy	Hanks
...		

- The *index selection problem*
 - Given a table, and a “workload” (big Java application with lots of SQL queries), decide which indexes to create (and which ones NOT to create!)
- Who does index selection:
 - The database administrator DBA
 - Semi-automatically, using a database administration tool

Which Indexes?

ID	fName	IName
10	Tom	Hanks
20	Amy	Hanks
...		

- The *index selection problem*
 - Given a table, and a “workload” (big Java application with lots of SQL queries), decide which indexes to create (and which ones NOT to create!)
- Who does index selection:
 - The database administrator Did you know that the average DBA makes over \$100,000 per year? 
 - Semi-automatically, using a database administration tool

Index Selection: Which Search Key

- **Make some attribute K a search key if the WHERE clause contains:**
 - **An exact match on K**
 - **A range predicate on K**
 - **A join on K**

The Index Selection Problem 1

$V(M, N, P);$

Your workload is this

100000 queries:

```
SELECT *
FROM V
WHERE N=?
```

100 queries:

```
SELECT *
FROM V
WHERE P=?
```

The Index Selection Problem 1

V(M, N, P);

Your workload is this

100000 queries:

```
SELECT *
FROM V
WHERE N=?
```

100 queries:

```
SELECT *
FROM V
WHERE P=?
```

What indexes ?

The Index Selection Problem 1

$V(M, N, P);$

Your workload is this

100000 queries:

```
SELECT *
FROM V
WHERE N=?
```

100 queries:

```
SELECT *
FROM V
WHERE P=?
```

A: $V(N)$ and $V(P)$ (hash tables or B-trees)

The Index Selection Problem 2

$V(M, N, P);$

Your workload is this

100000 queries:

```
SELECT *  
FROM V  
WHERE N>? and N<?
```

100 queries:

```
SELECT *  
FROM V  
WHERE P=?
```

100000 queries:

```
INSERT INTO V  
VALUES (?, ?, ?)
```

What indexes ?

The Index Selection Problem 2

$V(M, N, P)$:

Your workload is this

100000 queries:

```
SELECT *  
FROM V  
WHERE N>? and N<?
```

100 queries:

```
SELECT *  
FROM V  
WHERE P=?
```

100000 queries:

```
INSERT INTO V  
VALUES (?, ?, ?)
```

A: definitely $V(N)$ (must B-tree); unsure about $V(P)$

The Index Selection Problem 3

V(M, N, P);

Your workload is this

100000 queries: 1000000 queries: 100000 queries:

`SELECT *
FROM V
WHERE N=?`

`SELECT *
FROM V
WHERE N=? and P>?`

`INSERT INTO V
VALUES (?, ?, ?)`

What indexes ?

The Index Selection Problem 3

$V(M, N, P);$

Your workload is this

100000 queries: 1000000 queries: 100000 queries:

`SELECT *
FROM V
WHERE N=?`

`SELECT *
FROM V
WHERE N=? and P>?`

`INSERT INTO V
VALUES (?, ?, ?)`

A: $V(N, P)$

How does this index differ from:
1. Two indexes $V(N)$ and $V(P)$?
2. An index $V(P, N)$?

The Index Selection Problem 4

V(M, N, P);

Your workload is this

1000 queries:

```
SELECT *  
FROM V  
WHERE N>? and N<?
```

100000 queries:

```
SELECT *  
FROM V  
WHERE P>? and P<?
```

What indexes ?

The Index Selection Problem 4

V(M, N, P);

Your workload is this

1000 queries:

```
SELECT *
FROM V
WHERE N>? and N<?
```

100000 queries:

```
SELECT *
FROM V
WHERE P>? and P<?
```

A: V(N) secondary, V(P) primary index

Two typical kinds of queries

```
SELECT *
FROM Movie
WHERE year = ?
```

```
SELECT *
FROM Movie
WHERE year >= ? AND
      year <= ?
```

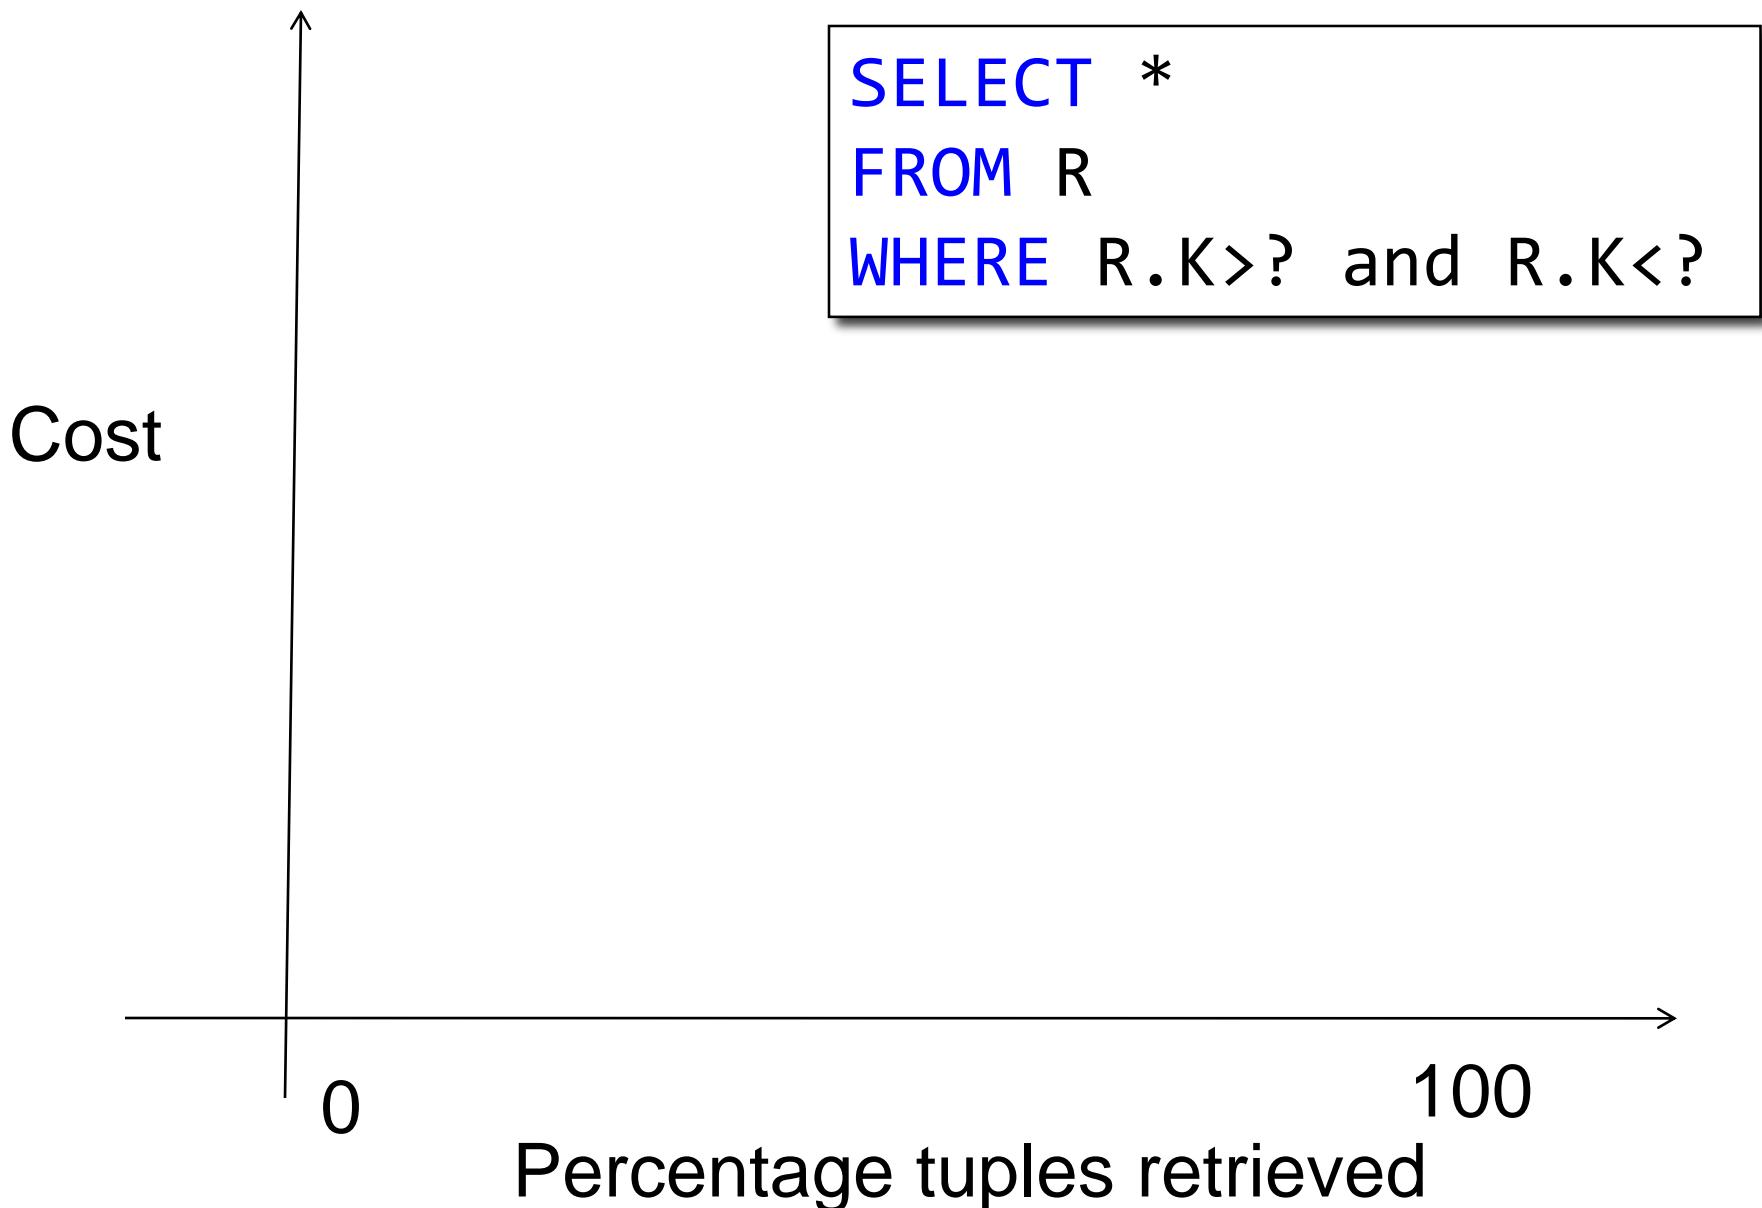
- **Point queries**
- **What data structure should be used for index?**
- **Range queries**
- **What data structure should be used for index?**

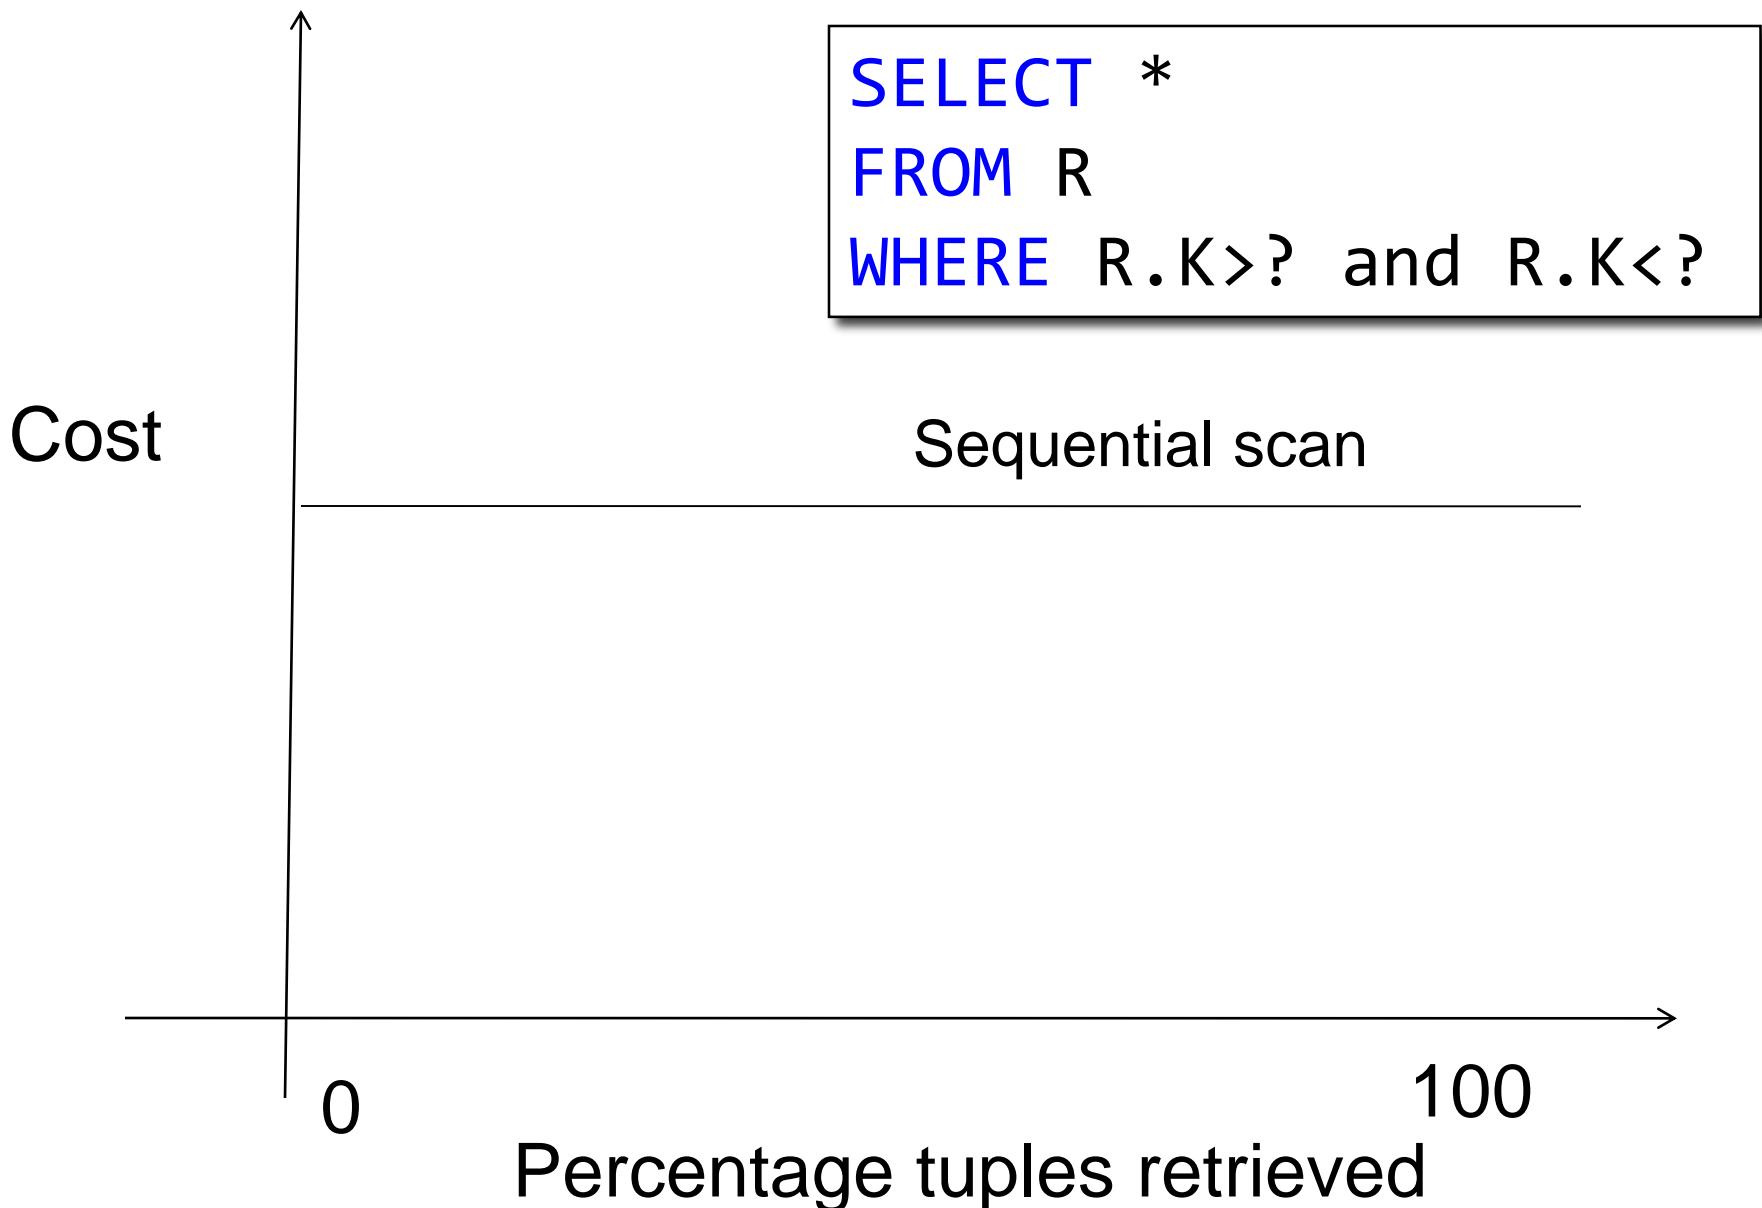
Basic Index Selection Guidelines

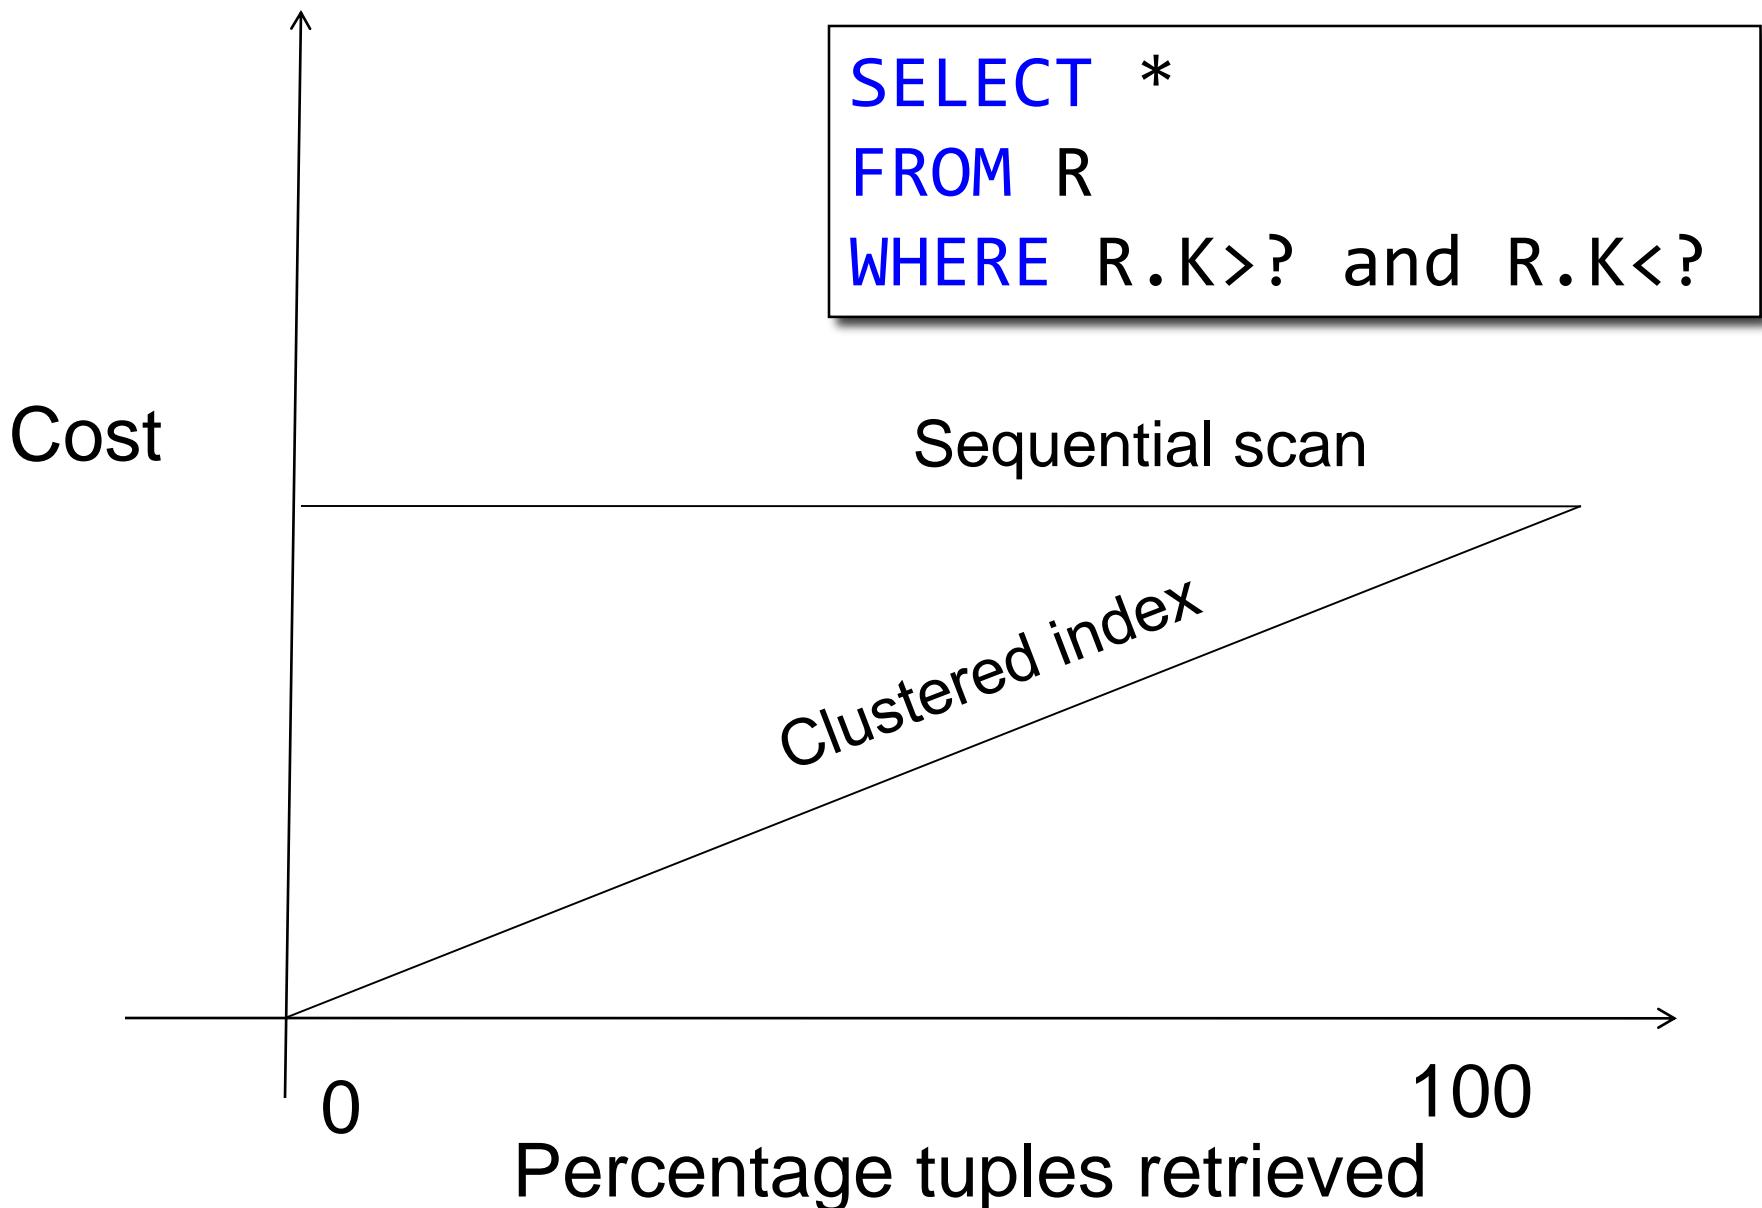
- **Consider queries in workload in order of importance**
- **Consider relations accessed by query**
 - No point indexing other relations
- **Look at WHERE clause for possible search key**
- **Try to choose indexes that speed-up multiple queries**

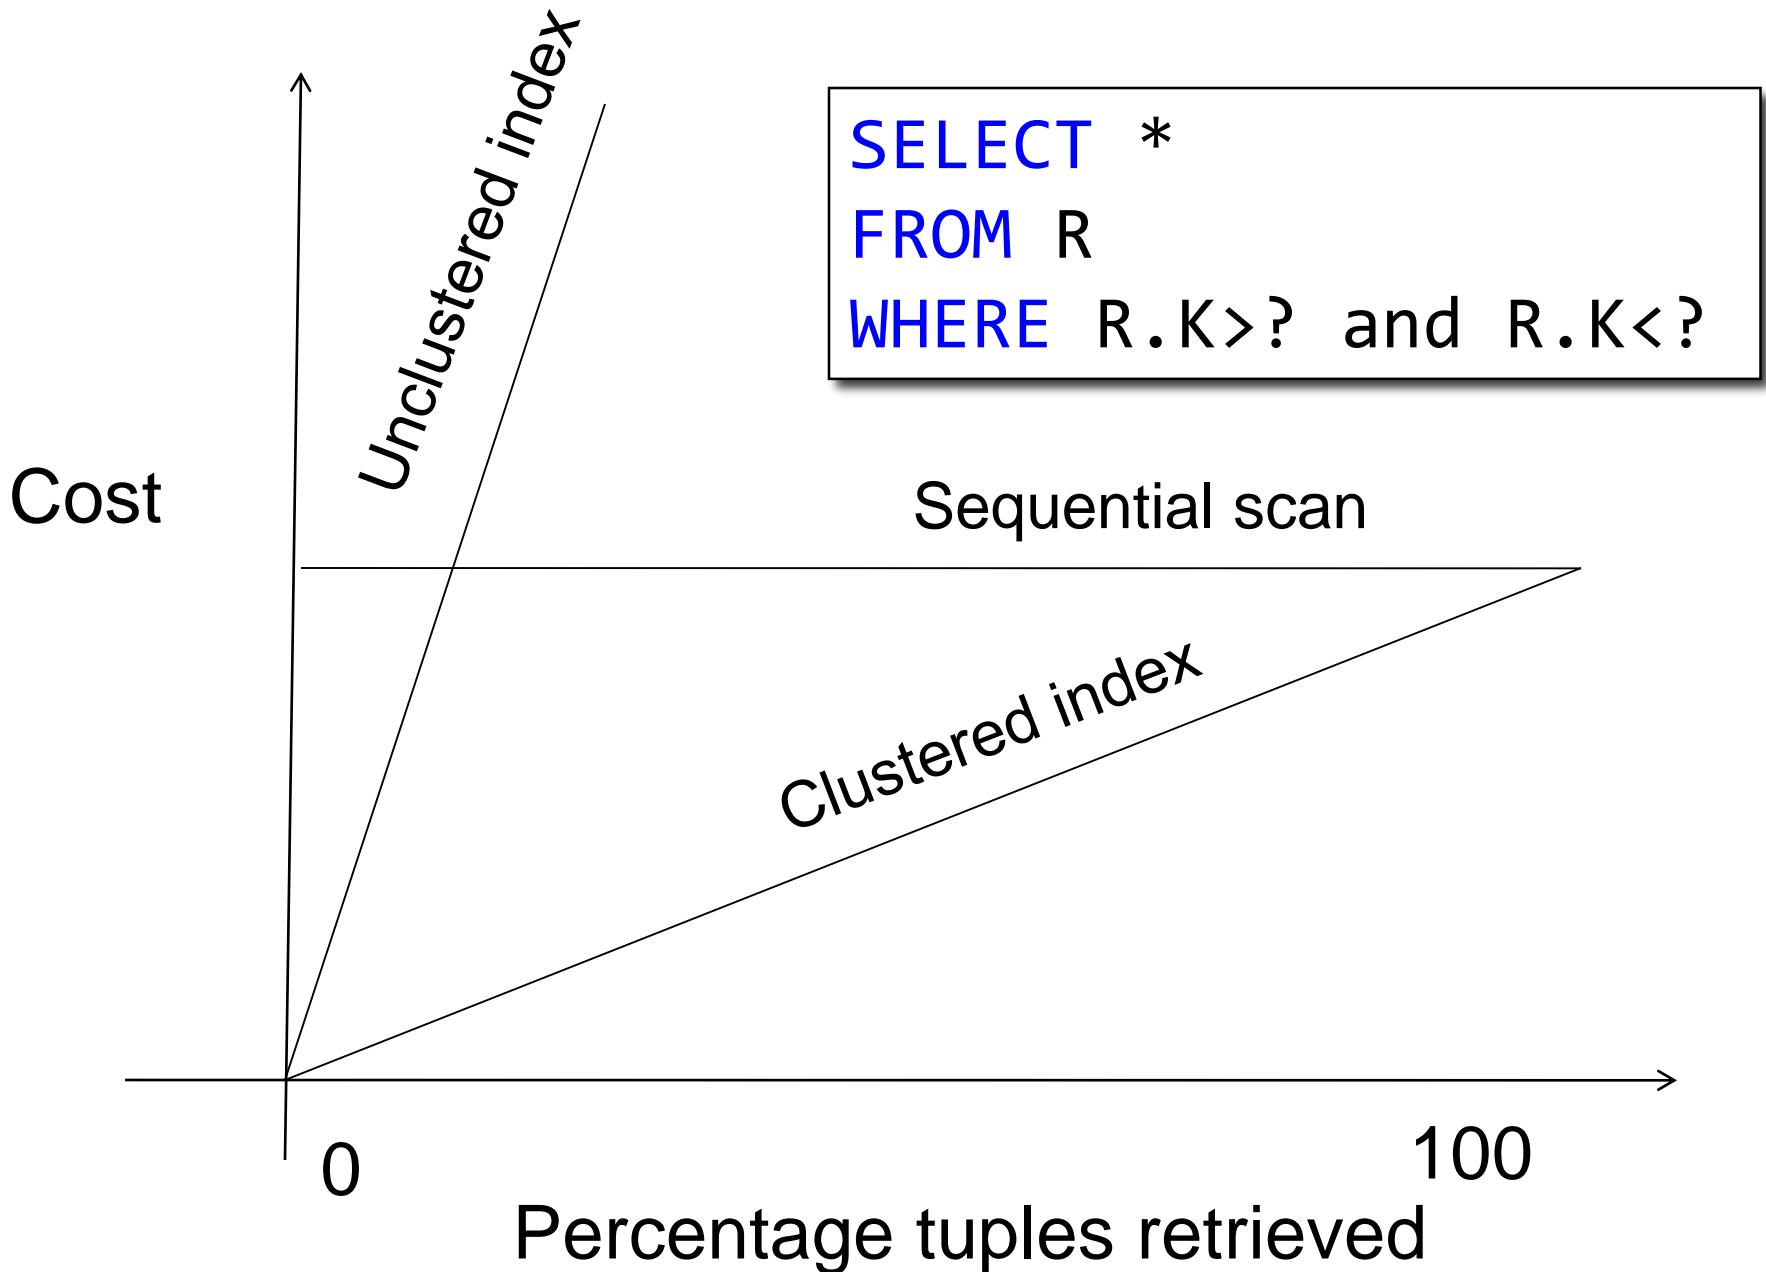
To Cluster or Not

- Range queries benefit mostly from clustering
- Covering indexes do *not* need to be clustered: they work equally well unclustered









Introduction to Database Systems

DATA 514

Lecture 17:
**Basics of Query Optimization and Query
Cost Estimation**

Choosing Index is Not Enough

- To estimate the cost of a query plan, we still need to consider other factors:
 - How each operator is implemented
 - The cost of each operator
 - Let's start with the basics

Cost of Reading Data From Disk

Cost Parameters

- **Cost = I/O + CPU + Network BW**
 - We will focus on I/O in this class
- **Parameters (a.k.a. statistics):**
 - $B(R)$ = # of blocks (i.e., pages) for relation R
 - $T(R)$ = # of tuples in relation R
 - $V(R, a)$ = # of distinct values of attribute a

Cost Parameters

- **Cost = I/O + CPU + Network BW**
 - We will focus on I/O in this class
- **Parameters (a.k.a. statistics):**
 - **B(R) = # of blocks (i.e., pages) for relation R**
 - **T(R) = # of tuples in relation R**
 - **V(R, a) = # of distinct values of attribute a**

When **a** is a key, **V(R,a) = T(R)**

When **a** is not a key, **V(R,a)** can be anything <= **T(R)**

Cost Parameters

- **Cost = I/O + CPU + Network BW**
 - We will focus on I/O in this class
- **Parameters (a.k.a. statistics):**
 - **B(R) = # of blocks (i.e., pages) for relation R**
 - **T(R) = # of tuples in relation R**
 - **V(R, a) = # of distinct values of attribute a**

When **a** is a key, **V(R,a) = T(R)**

When **a** is not a key, **V(R,a)** can be anything <= **T(R)**

- **DBMS collects statistics about base tables must infer them for intermediate results**

Selectivity Factors for Conditions

- $A = c$ /* $\sigma_{A=c}(R)$ */
 - Selectivity = $1/V(R, A)$
- $A < c$ /* $\sigma_{A < c}(R)$ */
 - Selectivity = $(c - \min(R, A)) / (\max(R, A) - \min(R, A))$
- $c_1 < A < c_2$ /* $\sigma_{c_1 < A < c_2}(R)$ */
 - Selectivity = $(c_2 - c_1) / (\max(R, A) - \min(R, A))$

Cost of Reading Data From Disk

- Sequential scan for relation R costs $B(R)$
- Index-based selection
 - Estimate selectivity factor f (see previous slide)
 - Clustered index: $f^*B(R)$
 - Unclustered index $f^*T(R)$

Note: we ignore I/O cost for index pages

Index Based Selection

- **Example:**

$$\begin{aligned}B(R) &= 2000 \\T(R) &= 100,000 \\V(R, a) &= 20\end{aligned}$$

$$\text{cost of } \sigma_{a=v}(R) = ?$$

- **Table scan:**
- **Index based selection:**

Index Based Selection

- **Example:**

$$\begin{aligned}B(R) &= 2000 \\T(R) &= 100,000 \\V(R, a) &= 20\end{aligned}$$

$$\text{cost of } \sigma_{a=v}(R) = ?$$

- **Table scan: $B(R) = 2,000$ I/Os**
- **Index based selection:**

Index Based Selection

- **Example:**

$B(R) = 2000$
$T(R) = 100,000$
$V(R, a) = 20$

cost of $\sigma_{a=v}(R) = ?$

- **Table scan: $B(R) = 2,000 \text{ I/Os}$**
- **Index based selection:**
 - **If index is clustered:**
 - **If index is unclustered:**

Index Based Selection

- **Example:**

$$\begin{aligned}B(R) &= 2000 \\T(R) &= 100,000 \\V(R, a) &= 20\end{aligned}$$

$$\text{cost of } \sigma_{a=v}(R) = ?$$

- **Table scan: $B(R) = 2,000$ I/Os**
- **Index based selection:**
 - If index is clustered: $B(R) * 1/V(R,a) = 100$ I/Os
 - If index is unclustered:

Index Based Selection

- **Example:**

$$\begin{aligned}B(R) &= 2000 \\T(R) &= 100,000 \\V(R, a) &= 20\end{aligned}$$

$$\text{cost of } \sigma_{a=v}(R) = ?$$

- **Table scan: $B(R) = 2,000$ I/Os**
- **Index based selection:**
 - If index is clustered: $B(R) * 1/V(R,a) = 100$ I/Os
 - If index is unclustered: $T(R) * 1/V(R,a) = 5,000$ I/Os

Index Based Selection

- Example:

$$\begin{aligned}B(R) &= 2000 \\T(R) &= 100,000 \\V(R, a) &= 20\end{aligned}$$

$$\text{cost of } \sigma_{a=v}(R) = ?$$

- Table scan: $B(R) = 2,000 \text{ I/Os}$
- Index based selection:
 - If index is clustered: $B(R) * 1/V(R,a) = 100 \text{ I/Os}$
 - If index is unclustered: $T(R) * 1/V(R,a) = 5,000 \text{ I/Os}$

Lesson: Don't build unclustered indexes when $V(R,a)$ is small !

Cost of Executing Operators (Focus on Joins)

Outline

- **Join operator algorithms**
 - One-pass algorithms (Sec. 15.2 and 15.3)
 - Index-based algorithms (Sec 15.6)
- **Note about readings:**
 - In class, we discuss only algorithms for joins
 - Other operators are easier: read the book

Join Algorithms

- Hash join
- Nested loop join
- Sort-merge join

Hash Join

Hash join: $R \bowtie S$

- Scan R, build buckets in main memory
- Then scan S and join
- Cost: $B(R) + B(S)$
- Which relation to build the hash table on?

Hash Join

Hash join: $R \bowtie S$

- Scan R, build buckets in main memory
- Then scan S and join
- Cost: $B(R) + B(S)$
- Which relation to build the hash table on?
- One-pass algorithm when $B(R) \leq M$
 - M = number of memory pages available

Hash Join Example

Patient(pid, name, address)

Insurance(pid, provider, policy_nb)

Patient \bowtie Insurance

Two tuples per page

Patient

1	'Bob'	'Seattle'
2	'Ela'	'Everett'

3	'Jill'	'Kent'
4	'Joe'	'Seattle'

Insurance

2	'Blue'	123
4	'Prem'	432

4	'Prem'	343
3	'GrpH'	554

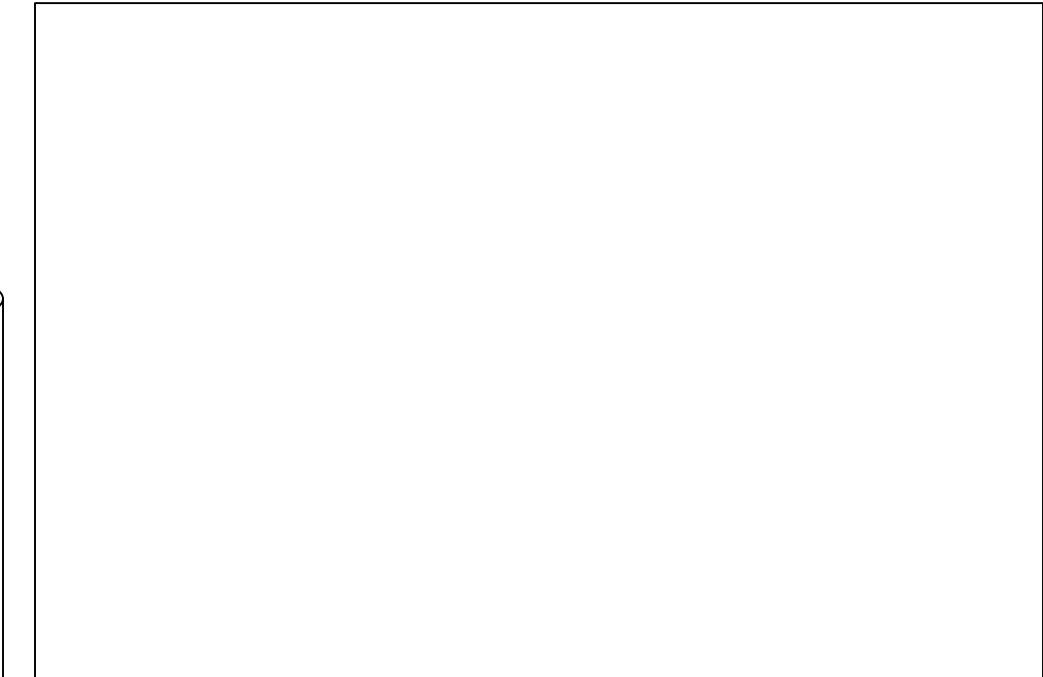
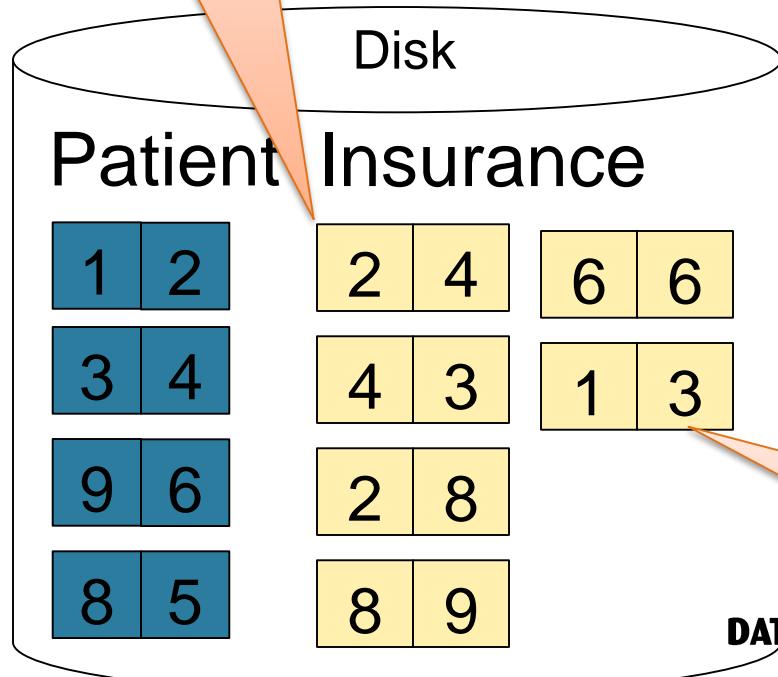
Hash Join Example

Patient \bowtie Insurance

Some large-enough #

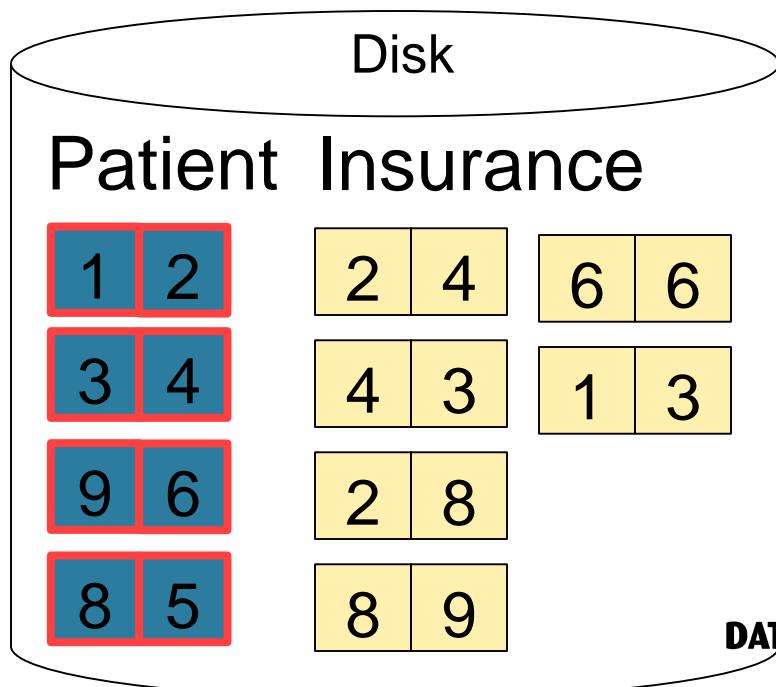
Memory M = 21 pages

Showing pid only



Hash Join Example

Step 1: Scan Patient and **build** hash table in memory
Can be done in
method open()



Memory M = 21 pages

Hash h: pid % 5

5		1	6	2		3	8	4	9
---	--	---	---	---	--	---	---	---	---

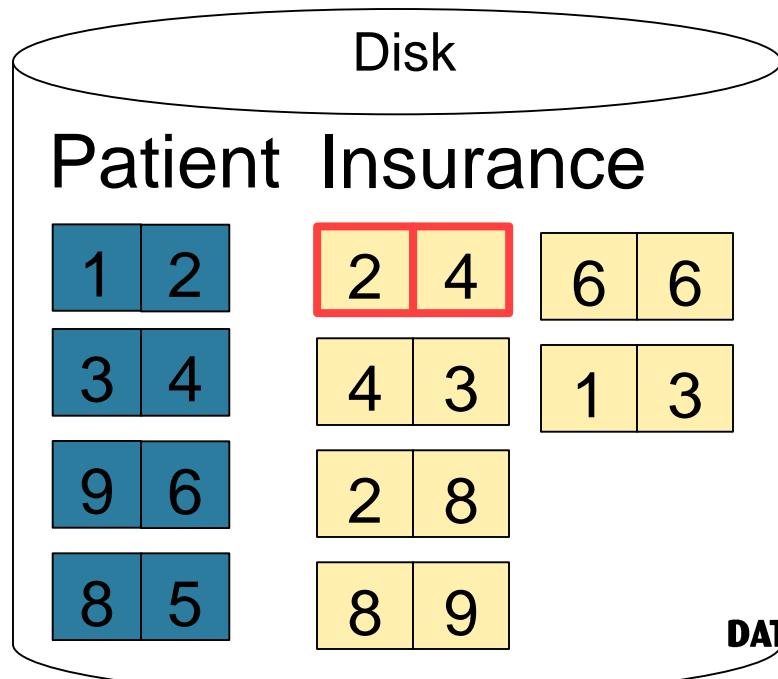


Input buffer

Hash Join Example

Step 2: Scan Insurance and **probe** into hash table

Done during
calls to next()



Memory M = 21 pages

Hash h: pid % 5

5		1	6	2		3	8	4	9
---	--	---	---	---	--	---	---	---	---

2	4
---	---

Input buffer

2	2
---	---

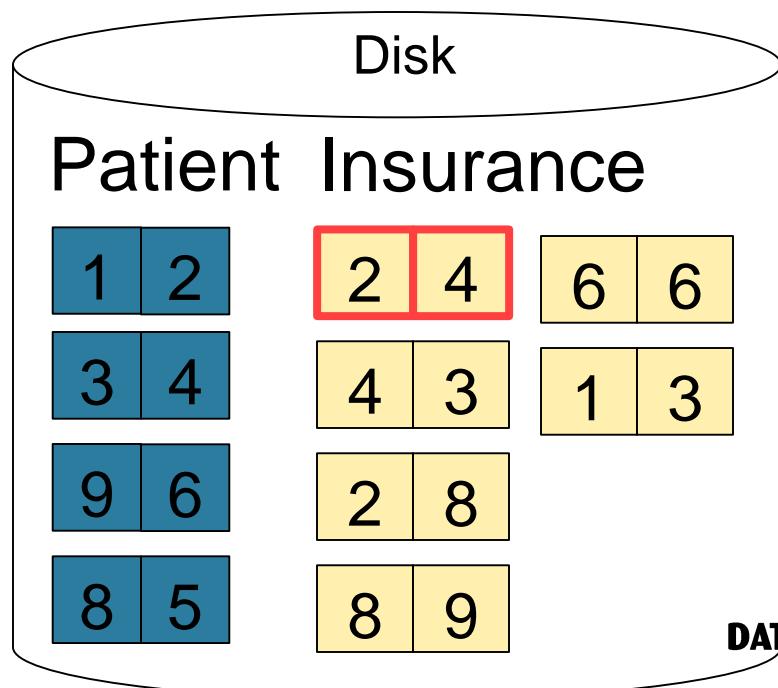
Output buffer

Write to disk or
pass to next
operator

Hash Join Example

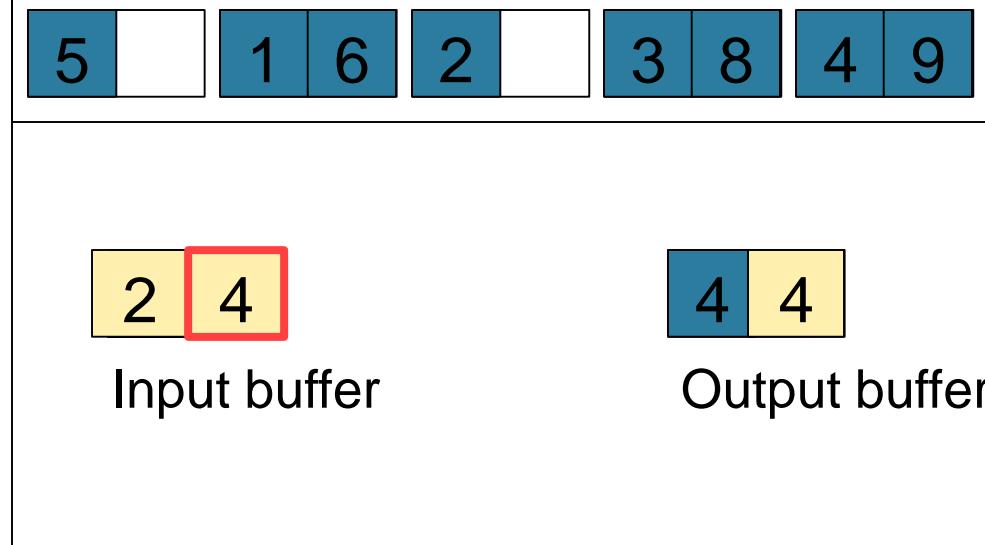
Step 2: Scan Insurance and **probe** into hash table

Done during
calls to next()



Memory M = 21 pages

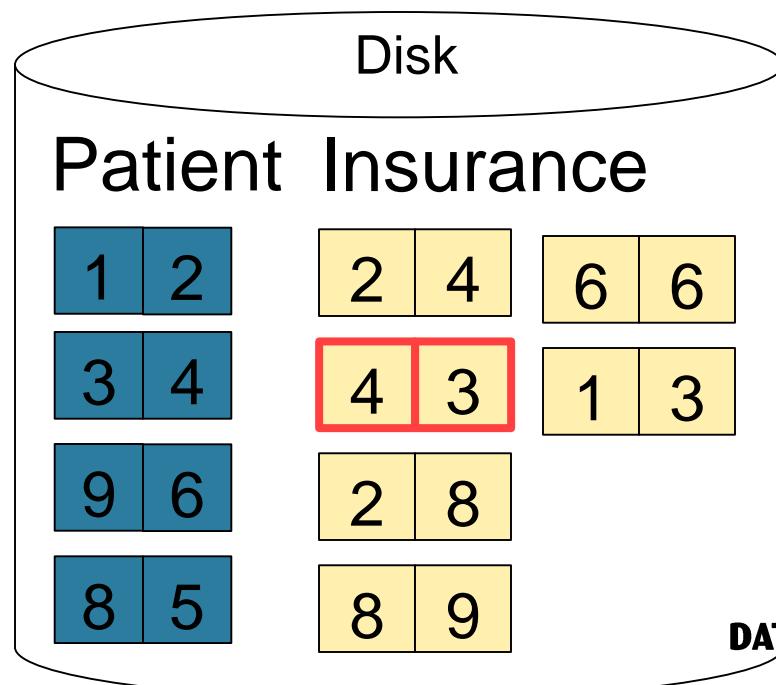
Hash h: pid % 5



Hash Join Example

Step 2: Scan Insurance and **probe** into hash table

Done during
calls to next()



Memory M = 21 pages

Hash h: pid % 5

5		1	6	2		3	8	4	9
---	--	---	---	---	--	---	---	---	---

4	3
---	---

Input buffer

4	4
---	---

Output buffer

Keep going until read all of Insurance

Cost: $B(R) + B(S)$

Nested Loop Joins

- **Tuple-based nested loop $R \bowtie S$**
- **R is the outer relation, S is the inner relation**

```
for each tuple t1 in R do  
  for each tuple t2 in S do  
    if t1 and t2 join then output (t1,t2)
```

What is the Cost?

Nested Loop Joins

- **Tuple-based nested loop $R \bowtie S$**
- **R is the outer relation, S is the inner relation**

```
for each tuple  $t_1$  in  $R$  do  
  for each tuple  $t_2$  in  $S$  do  
    if  $t_1$  and  $t_2$  join then output  $(t_1, t_2)$ 
```

What is the Cost?

- **Cost: $B(R) + T(R) B(S)$**
- **Multiple-pass since S is read many times**

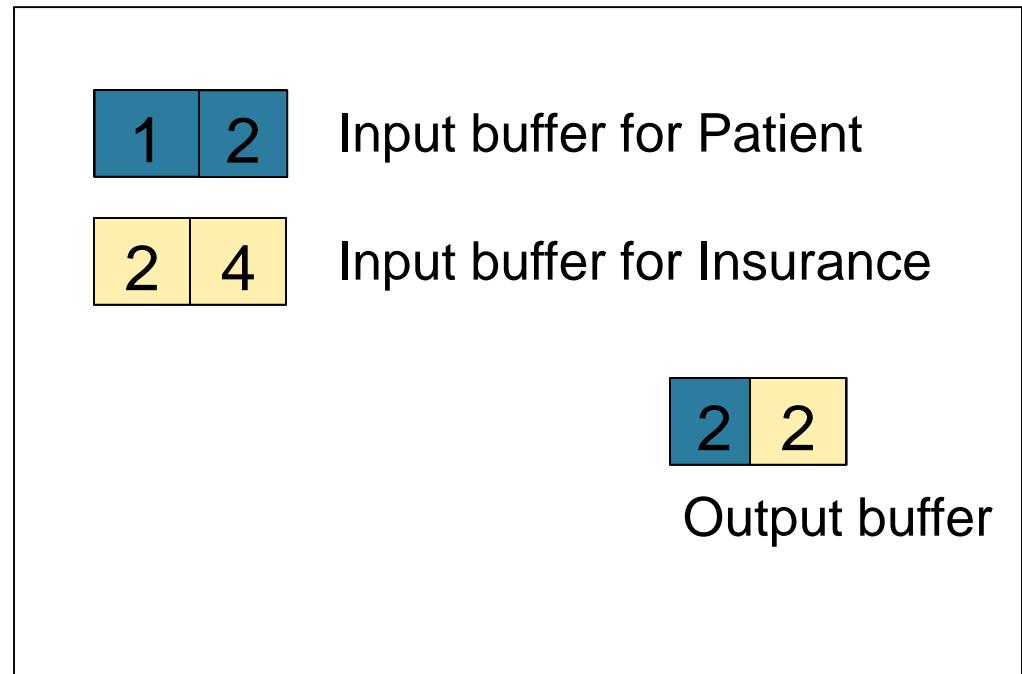
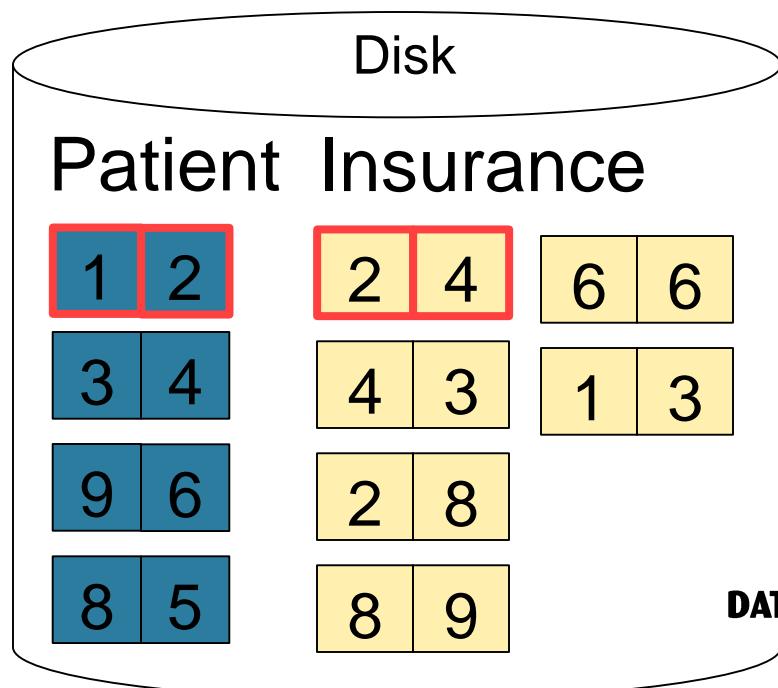
Page-at-a-time Refinement

```
for each page of tuples r in R do  
  for each page of tuples s in S do  
    for all pairs of tuples t1 in r, t2 in s  
      if t1 and t2 join then output (t1,t2)
```

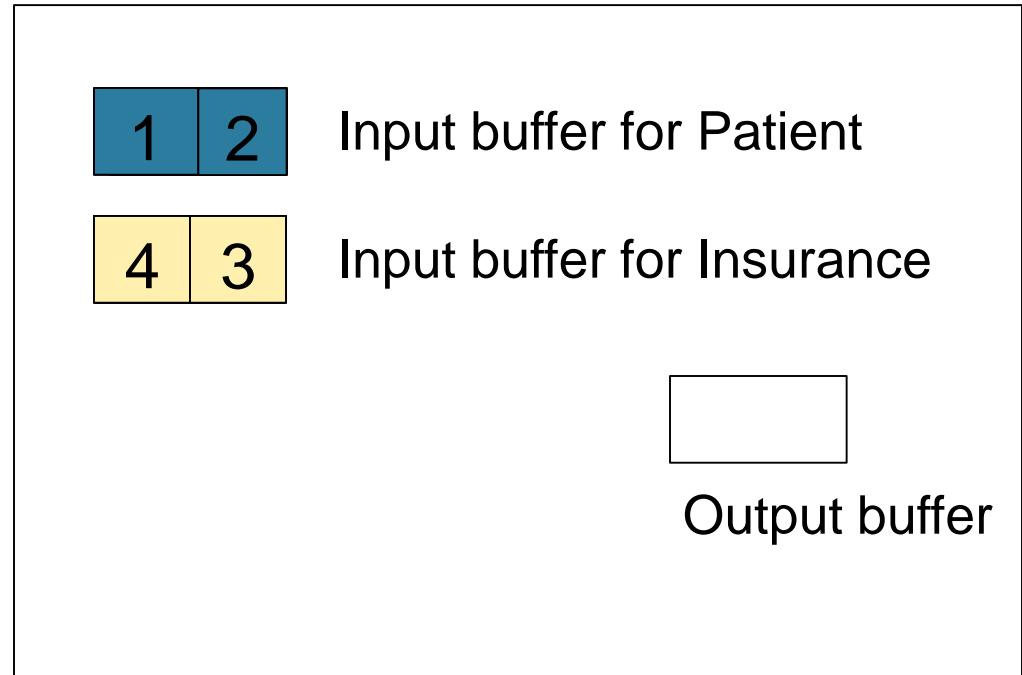
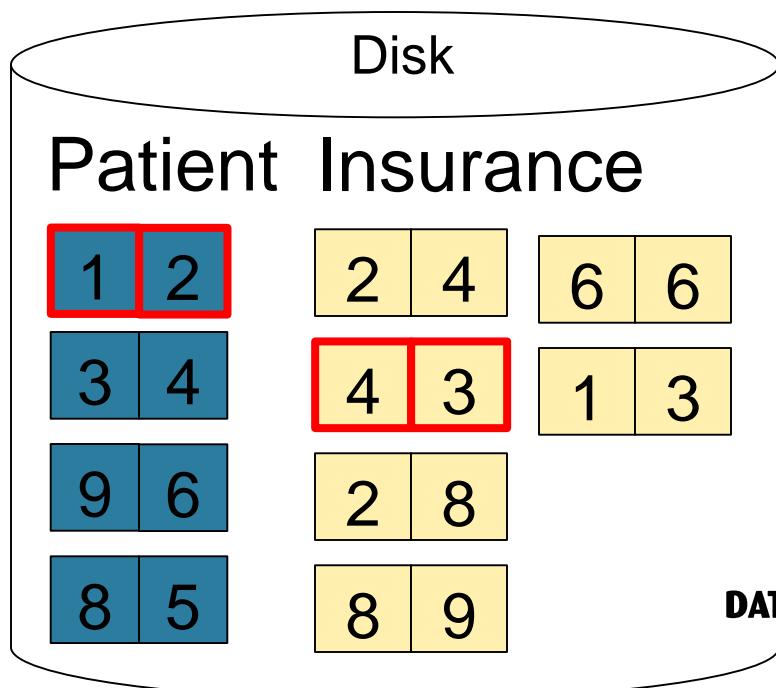
- Cost: $B(R) + B(R)B(S)$

What is the Cost?

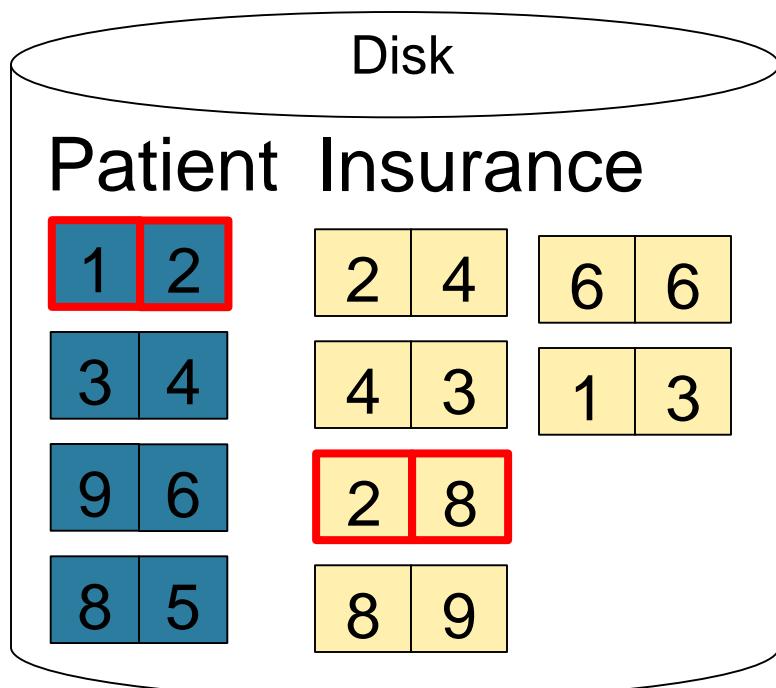
Page-at-a-time Refinement



Page-at-a-time Refinement



Page-at-a-time Refinement



1 | 2

Input buffer for Patient

2 | 8

Input buffer for Insurance

Keep going until read
all of Insurance

2 | 2

Then repeat for next
page of Patient... until end of Patient

Output buffer

Cost: $B(R) + B(R)B(S)$

Block-Nested-Loop Refinement

```
for each group of M-1 pages r in R do  
  for each page of tuples s in S do  
    for all pairs of tuples t1 in r, t2 in s  
      if t1 and t2 join then output (t1,t2)
```

- Cost: $B(R) + B(R)B(S)/(M-1)$

What is the Cost?

Sort-Merge Join

Sort-merge join: $R \bowtie S$

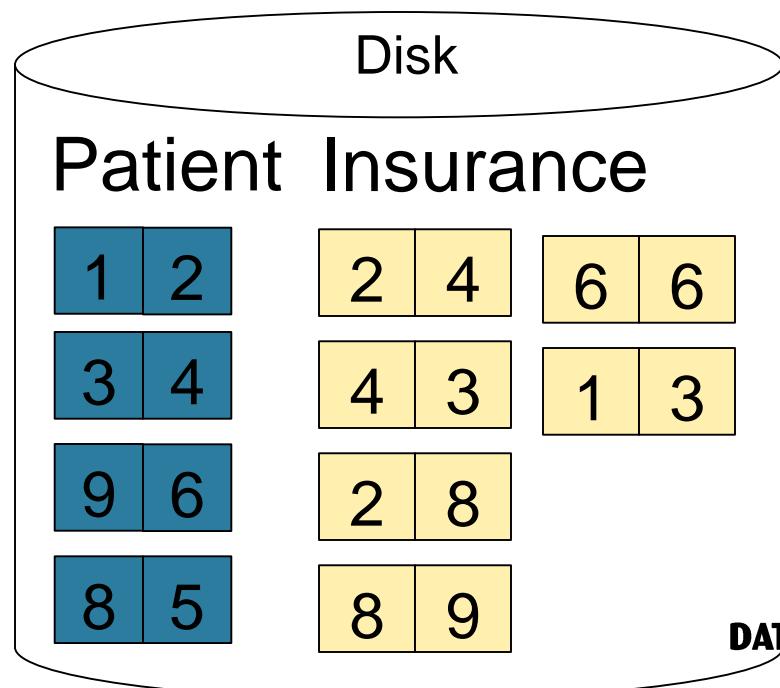
- **Scan R and sort in main memory**
 - **Scan S and sort in main memory**
 - **Merge R and S**
-
- **Cost: $B(R) + B(S)$**
 - **One pass algorithm when $B(S) + B(R) \leq M$**
 - **Typically, this is NOT a one pass algorithm**

Sort-Merge Join Example

Step 1: Scan Patient and **sort** in memory

Memory M = 21 pages

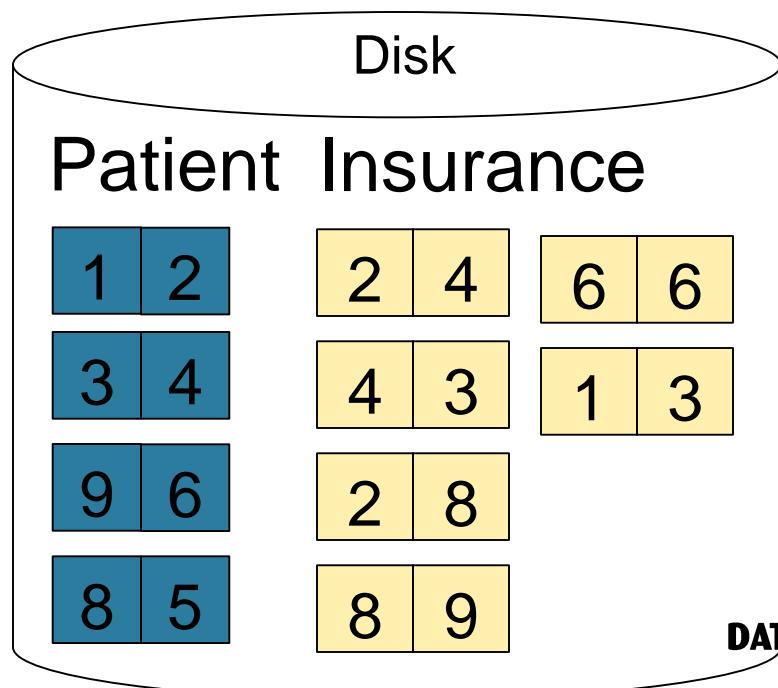
1	2	3	4	5	6	8	9
---	---	---	---	---	---	---	---



Sort-Merge Join Example

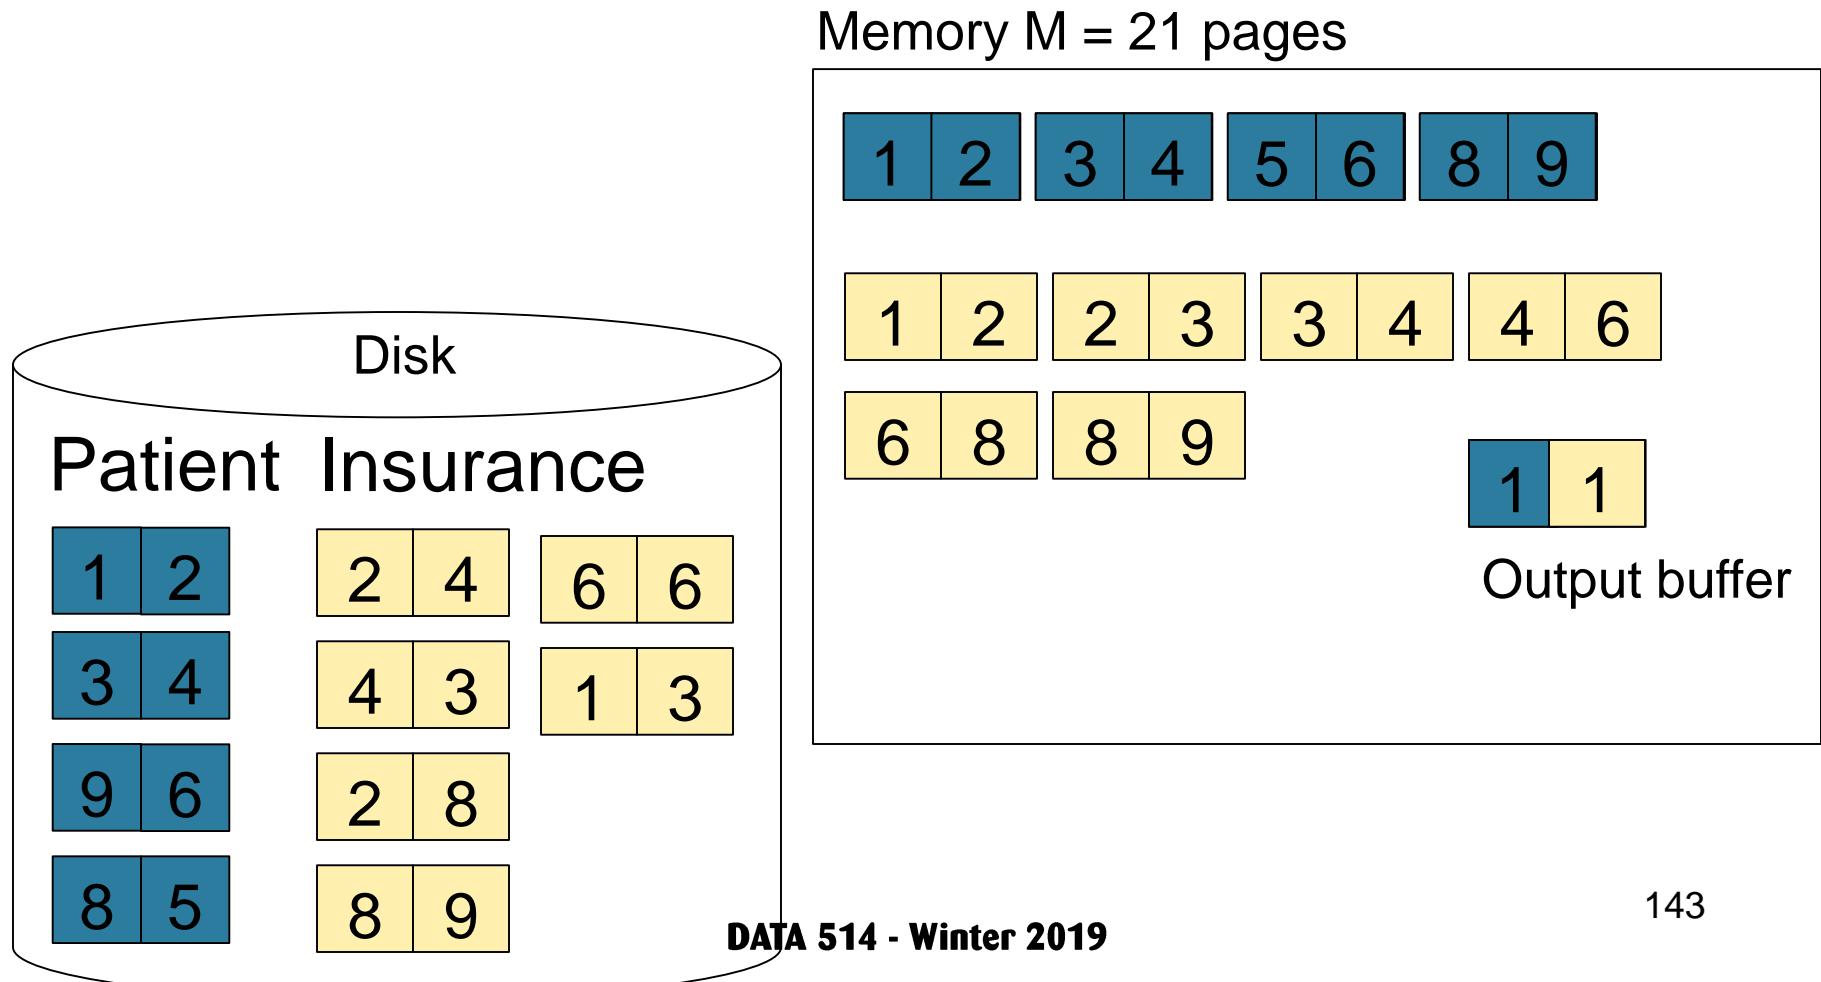
Step 2: Scan Insurance and **sort** in memory

Memory M = 21 pages



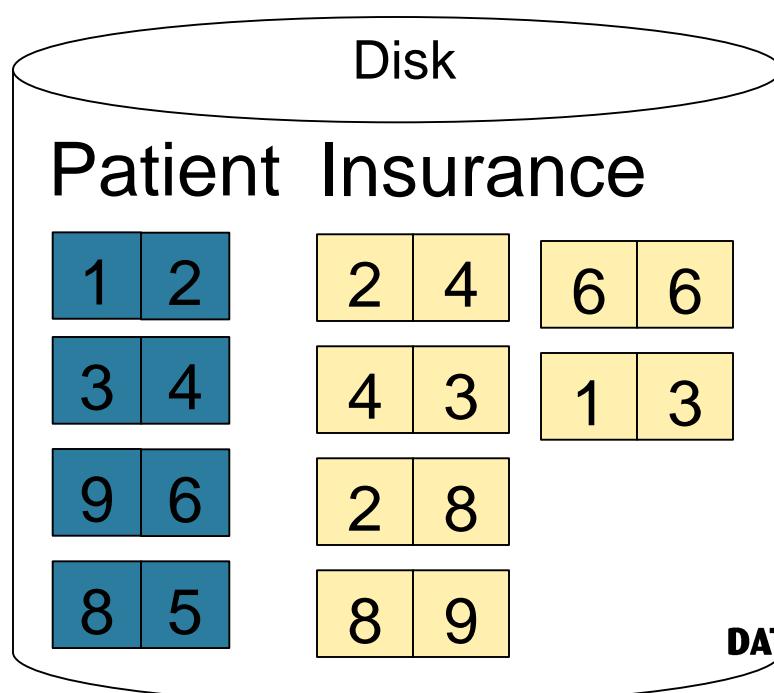
Sort-Merge Join Example

Step 3: Merge Patient and Insurance

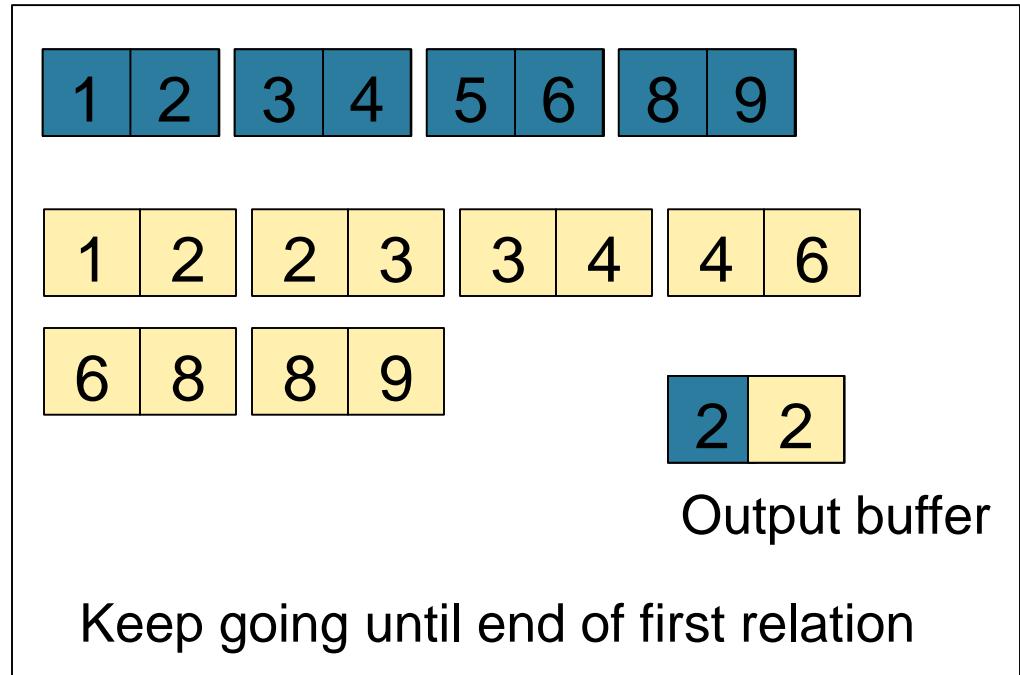


Sort-Merge Join Example

Step 3: Merge Patient and Insurance



Memory M = 21 pages



Index Nested Loop Join

$R \bowtie S$

- Assume S has an index on the join attribute
- Iterate over R , for each tuple fetch corresponding tuple(s) from S
- Cost:
 - If index on S is clustered:
$$B(R) + T(R) * (B(S) * 1/v(S,a))$$
 - If index on S is unclustered:
$$B(R) + T(R) * (f(S) * 1/v(S,a))$$

Cost of Query Plans

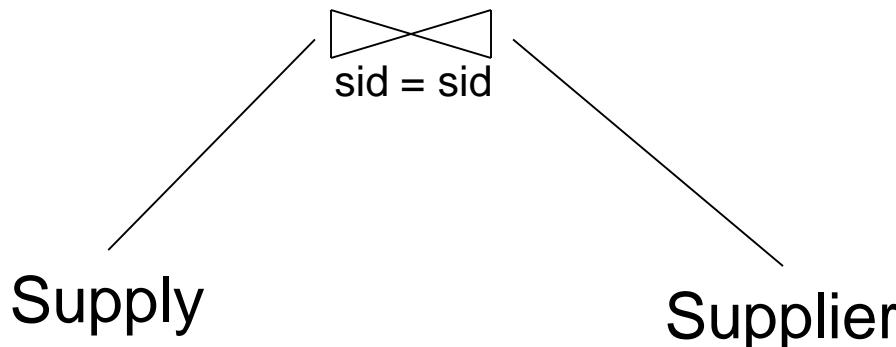
Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Logical Query Plan 1

π_{sname}

$\sigma_{\text{pno}=2 \wedge \text{scity}=\text{'Seattle'} \wedge \text{sstate}=\text{'WA'}}$



```
SELECT sname  
FROM Supplier x, Supply y  
WHERE x.sid = y.sid  
and y.pno = 2  
and x.scity = 'Seattle'  
and x.sstate = 'WA'
```

T(Supply) = 10000
B(Supply) = 100
V(Supply, pno) = 2500

T(Supplier) = 1000
B(Supplier) = 100
V(Supplier, scity) = 20
V(Supplier, sstate) = 10

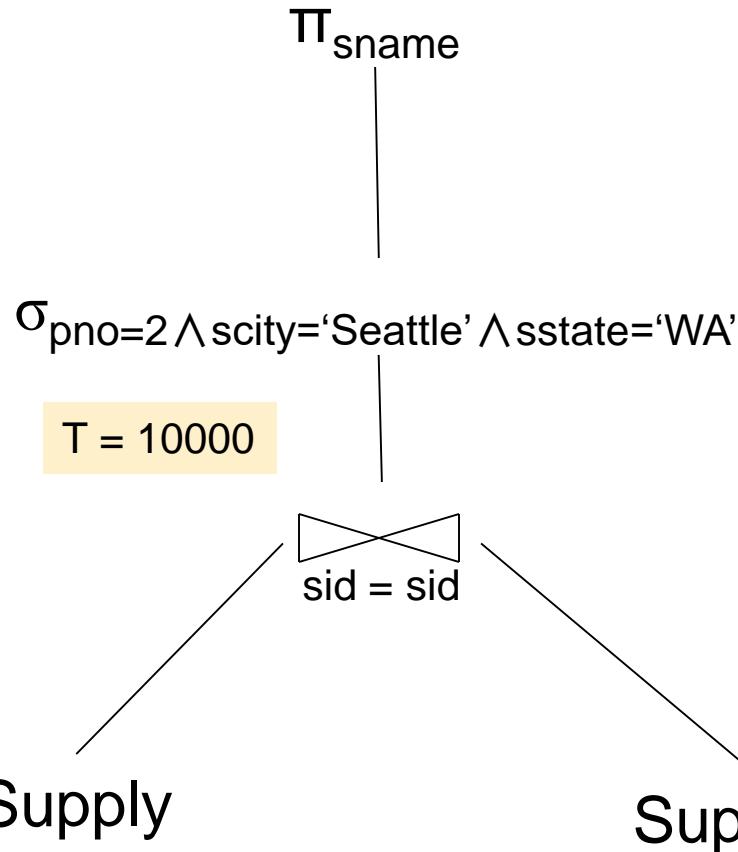
DATA 514 Winter 2019

M=11₁₄₇

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Logical Query Plan 1



```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
and y.pno = 2
and x.scity = 'Seattle'
and x.sstate = 'WA'
```

$T(\text{Supply}) = 10000$
 $B(\text{Supply}) = 100$
 $V(\text{Supply}, \text{pno}) = 2500$

$T(\text{Supplier}) = 1000$
 $B(\text{Supplier}) = 100$
 $V(\text{Supplier}, \text{scity}) = 20$
 $V(\text{Supplier}, \text{sstate}) = 10$

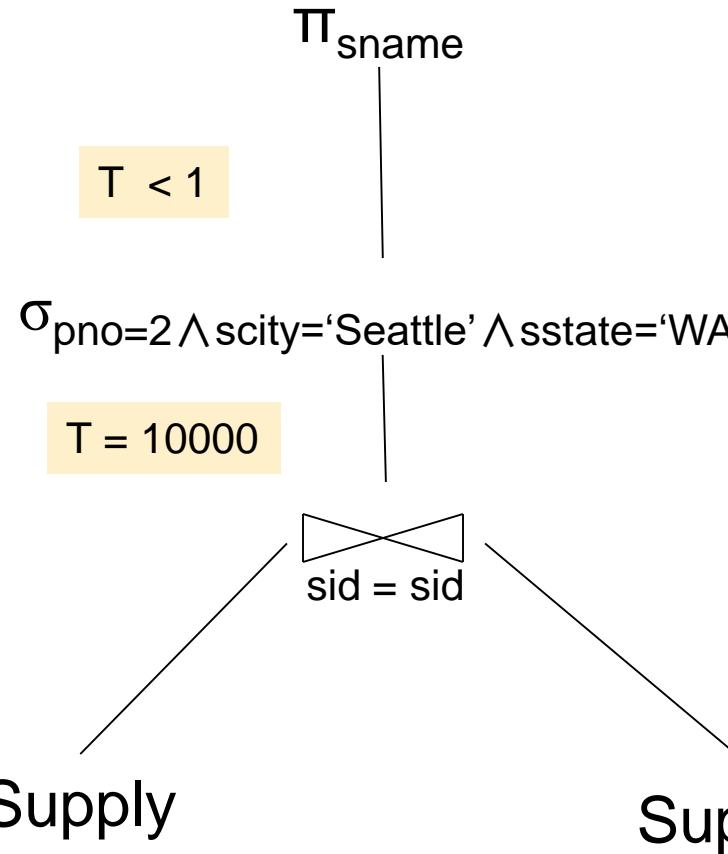
DATA 514 Winter 2019

M=11₁₄₈

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Logical Query Plan 1



```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
and y.pno = 2
and x.scity = 'Seattle'
and x.sstate = 'WA'
```

$T(\text{Supply}) = 10000$
 $B(\text{Supply}) = 100$
 $V(\text{Supply}, \text{pno}) = 2500$

$T(\text{Supplier}) = 1000$
 $B(\text{Supplier}) = 100$
 $V(\text{Supplier}, \text{scity}) = 20$
 $V(\text{Supplier}, \text{sstate}) = 10$

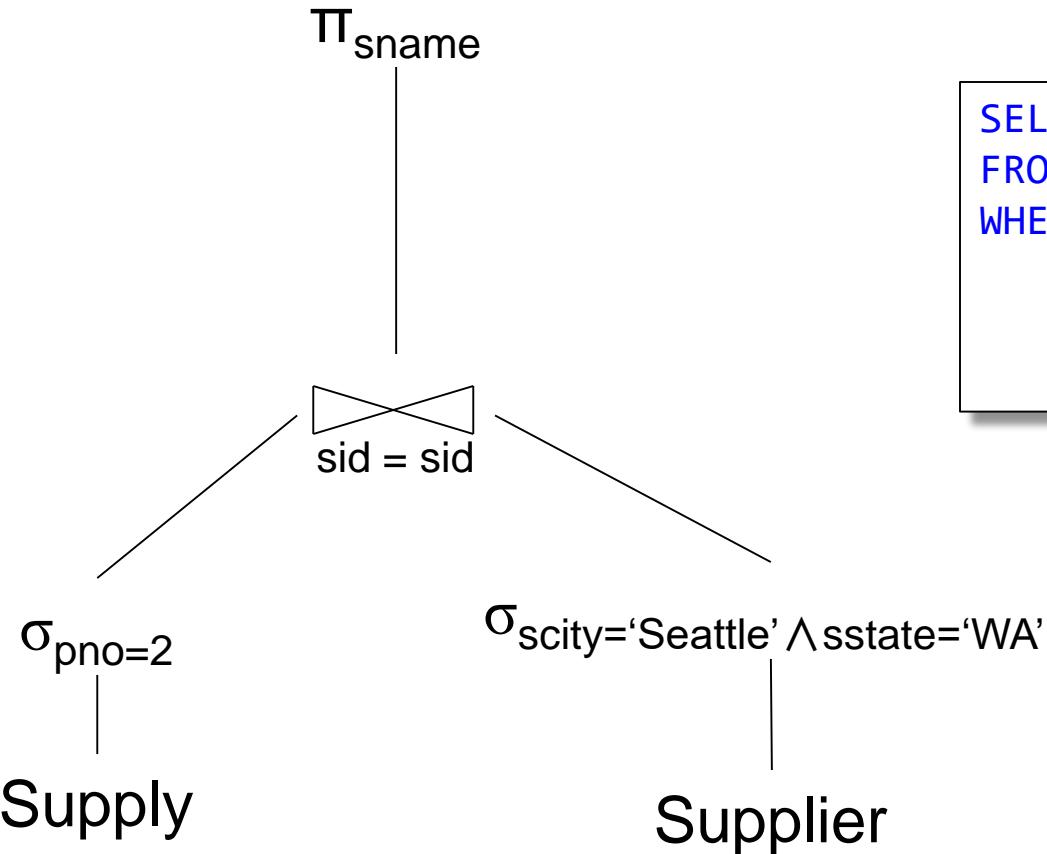
DATA 514 Winter 2019

M=11₁₄₉

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Logical Query Plan 2



```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
and y.pno = 2
and x.scity = 'Seattle'
and x.sstate = 'WA'
```

T(Supply) = 10000
B(Supply) = 100
V(Supply, pno) = 2500

T(Supplier) = 1000
B(Supplier) = 100
V(Supplier, scity) = 20
V(Supplier, state) = 10

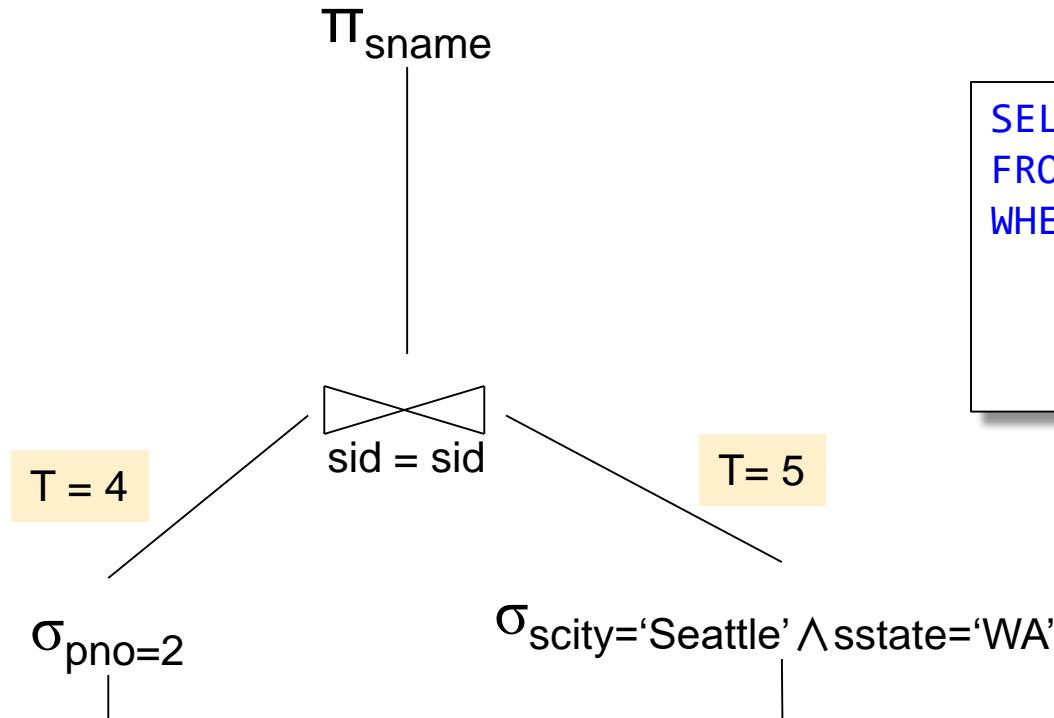
DATA 514 Winter 2019

M=11₁₅₀

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Logical Query Plan 2



```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
and y.pno = 2
and x.scity = 'Seattle'
and x.sstate = 'WA'
```

T(Supply) = 10000
B(Supply) = 100
V(Supply, pno) = 2500

T(Supplier) = 1000
B(Supplier) = 100
V(Supplier, scity) = 20
V(Supplier, state) = 10

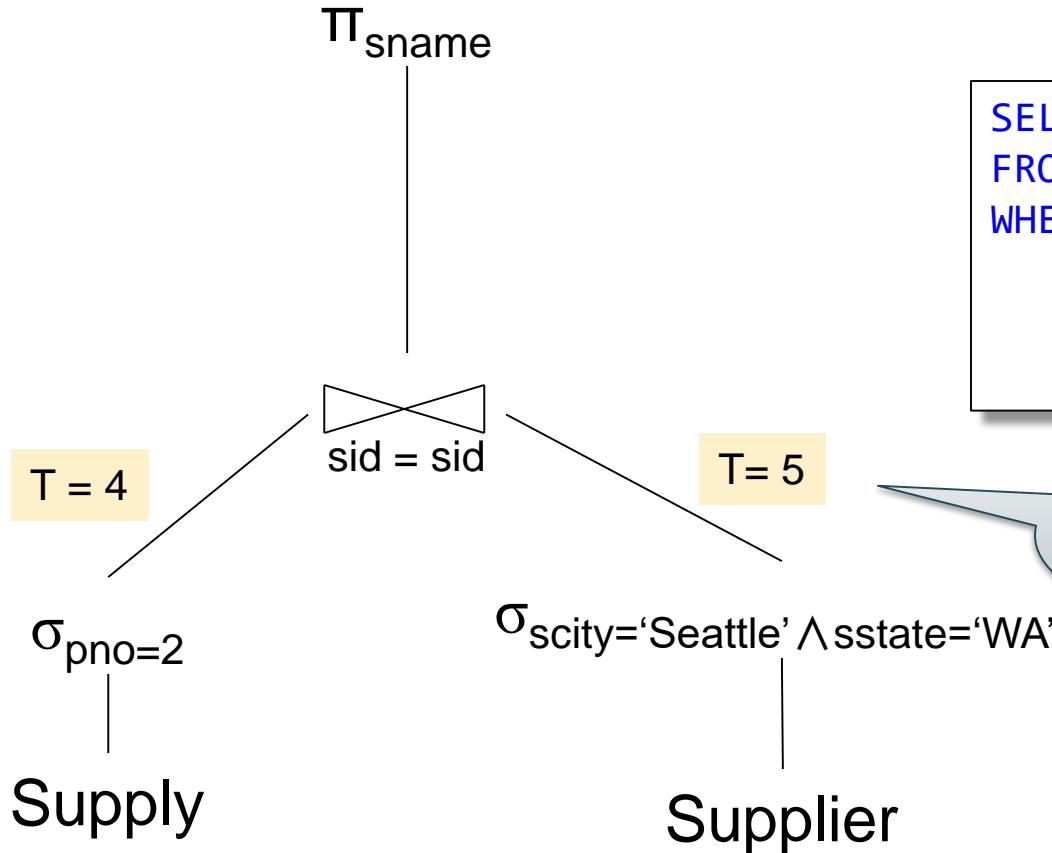
DATA 514 Winter 2019

M=11₁₅₁

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Logical Query Plan 2



```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
and y.pno = 2
and x.scity = 'Seattle'
and x.sstate = 'WA'
```

Very wrong!
Why?

T(Supply) = 10000
B(Supply) = 100
V(Supply, pno) = 2500

T(Supplier) = 1000
B(Supplier) = 100
V(Supplier, scity) = 20
V(Supplier, state) = 10

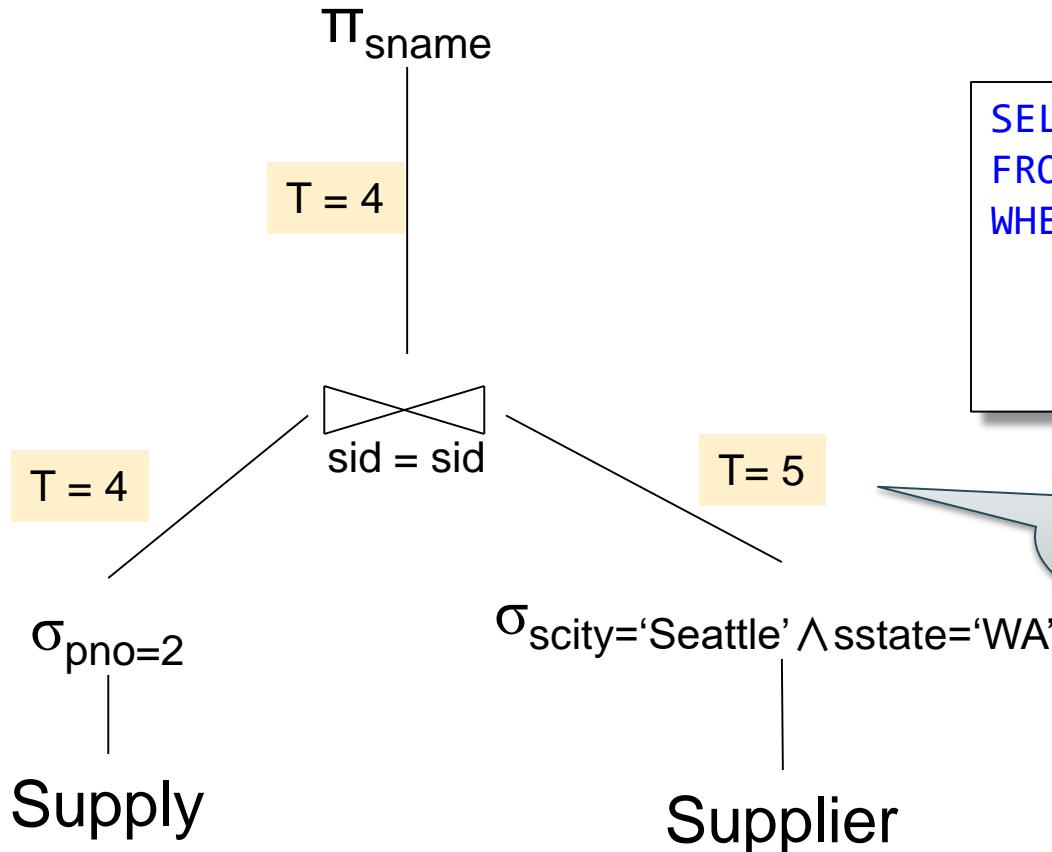
DATA 514 Winter 2019

M=11₁₅₂

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Logical Query Plan 2



Very wrong!
Why?

$T(\text{Supply}) = 10000$
 $B(\text{Supply}) = 100$
 $V(\text{Supply}, \text{pno}) = 2500$

$T(\text{Supplier}) = 1000$
 $B(\text{Supplier}) = 100$
 $V(\text{Supplier}, \text{scity}) = 20$
 $V(\text{Supplier}, \text{sstate}) = 10$

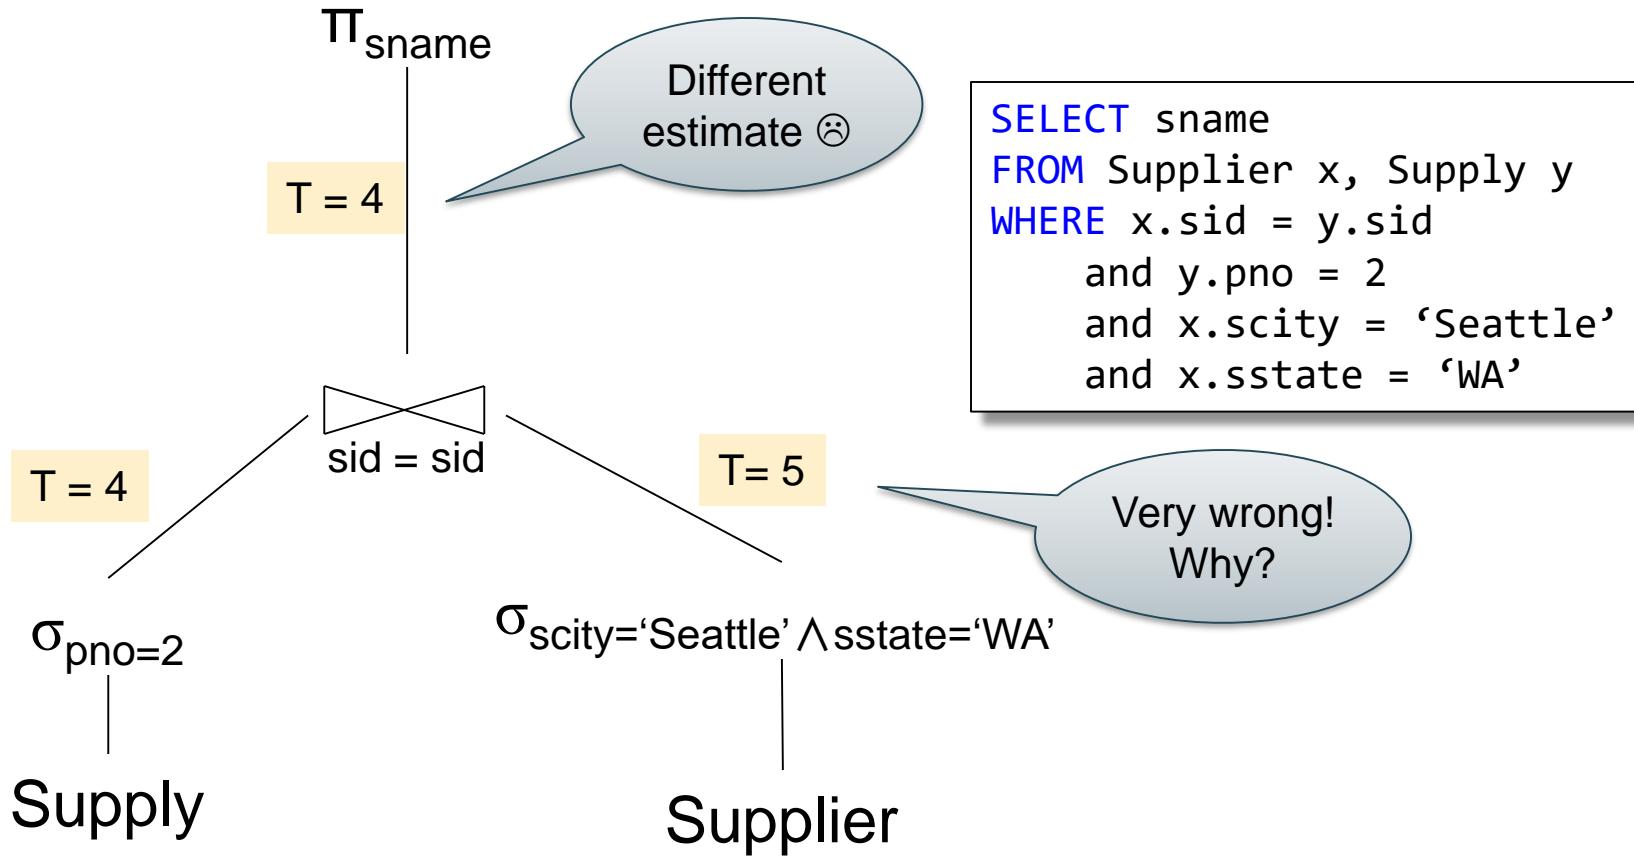
DATA 514 Winter 2019

M=11₁₅₃

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Logical Query Plan 2



$T(\text{Supply}) = 10000$
 $B(\text{Supply}) = 100$
 $V(\text{Supply}, \text{pno}) = 2500$

$T(\text{Supplier}) = 1000$
 $B(\text{Supplier}) = 100$
 $V(\text{Supplier}, \text{scity}) = 20$
 $V(\text{Supplier}, \text{sstate}) = 10$

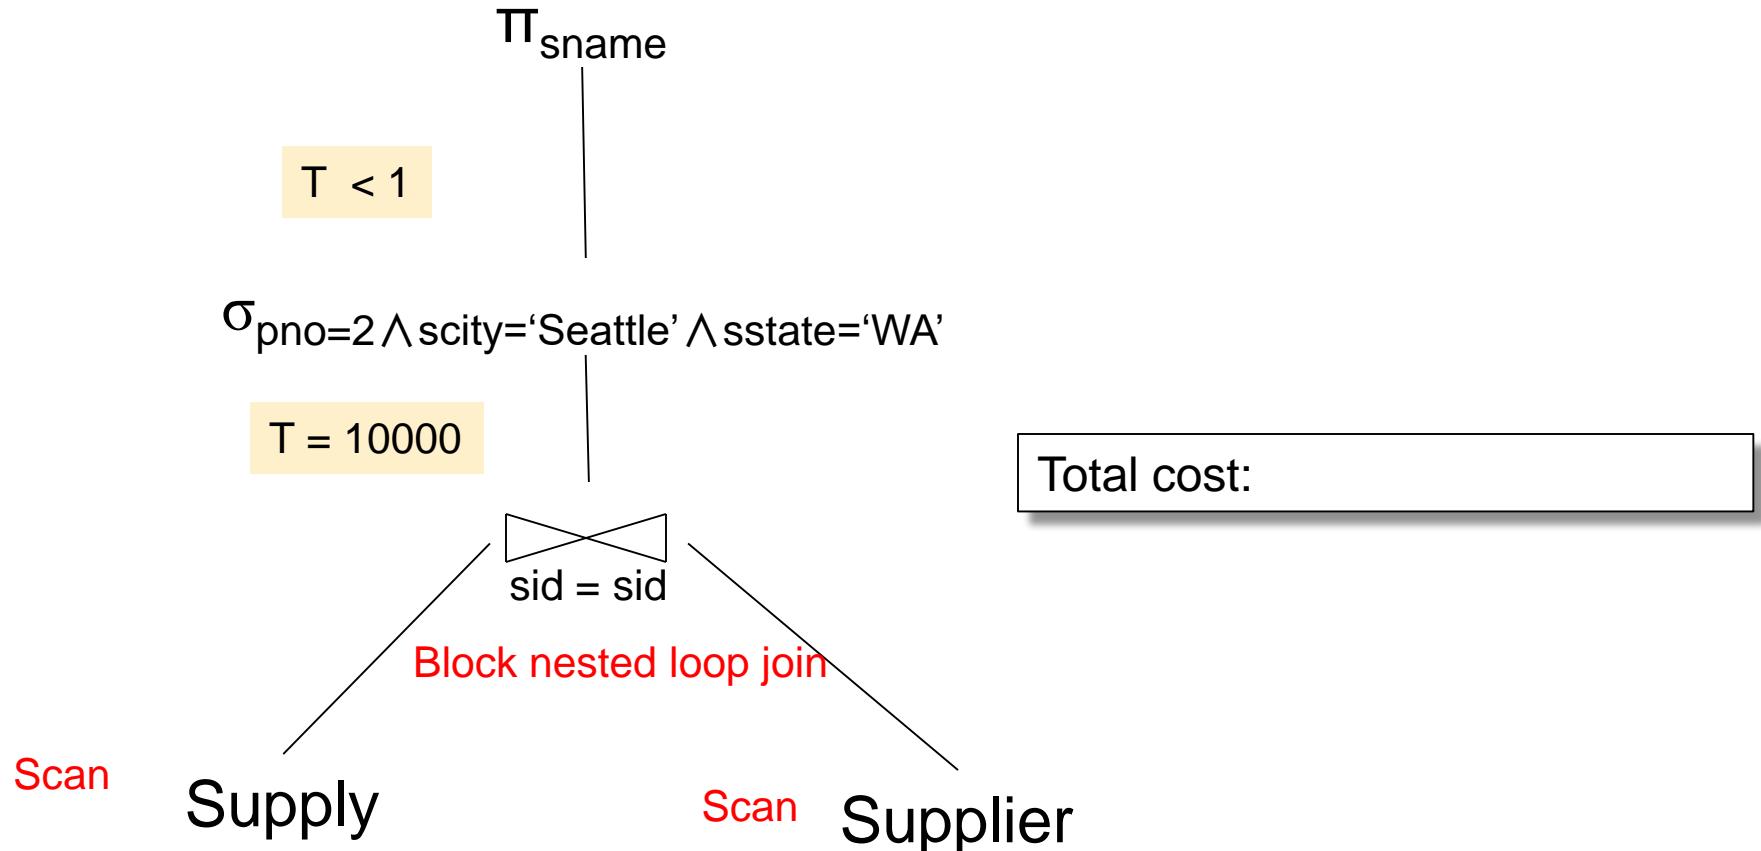
DATA 514 Winter 2019

M=11₁₅₄

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Physical Plan 1



$T(\text{Supply}) = 10000$
 $B(\text{Supply}) = 100$
 $V(\text{Supply}, \text{pno}) = 2500$

$T(\text{Supplier}) = 1000$
 $B(\text{Supplier}) = 100$
 $V(\text{Supplier}, \text{scity}) = 20$
 $V(\text{Supplier}, \text{sstate}) = 10$

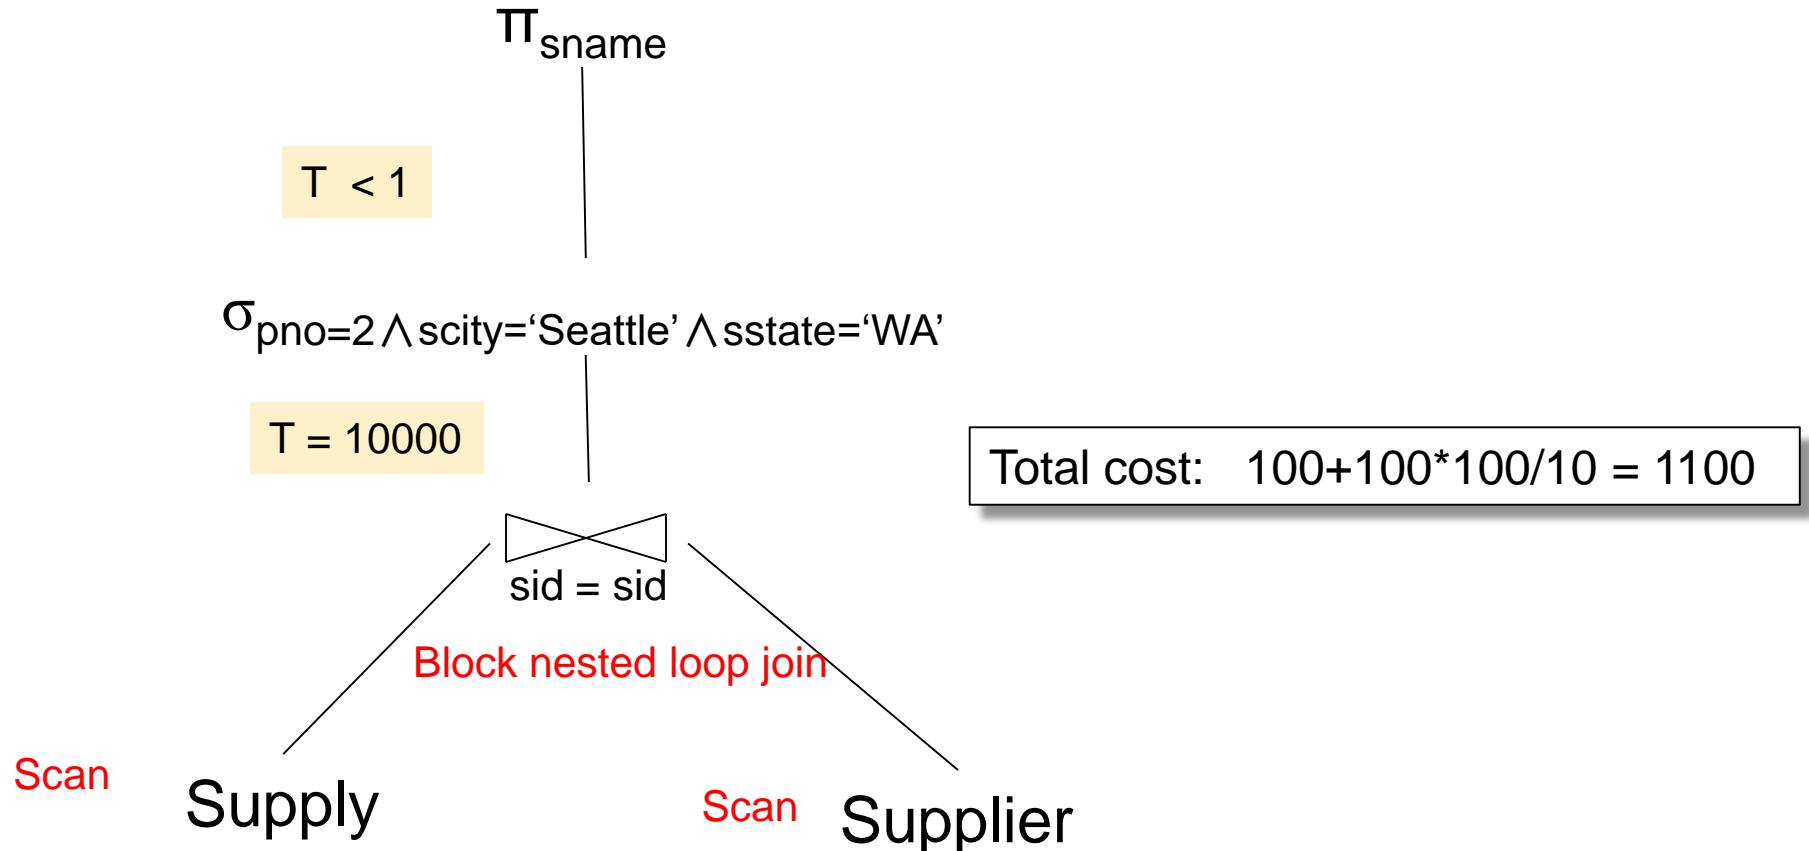
DATA 514 Winter 2019

M=11₁₅₅

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Physical Plan 1



$$T(\text{Supply}) = 10000$$

$$B(\text{Supply}) = 100$$

$$V(\text{Supply}, \text{pno}) = 2500$$

$$T(\text{Supplier}) = 1000$$

$$B(\text{Supplier}) = 100$$

$$V(\text{Supplier}, \text{scity}) = 20$$

$$V(\text{Supplier}, \text{sstate}) = 10$$

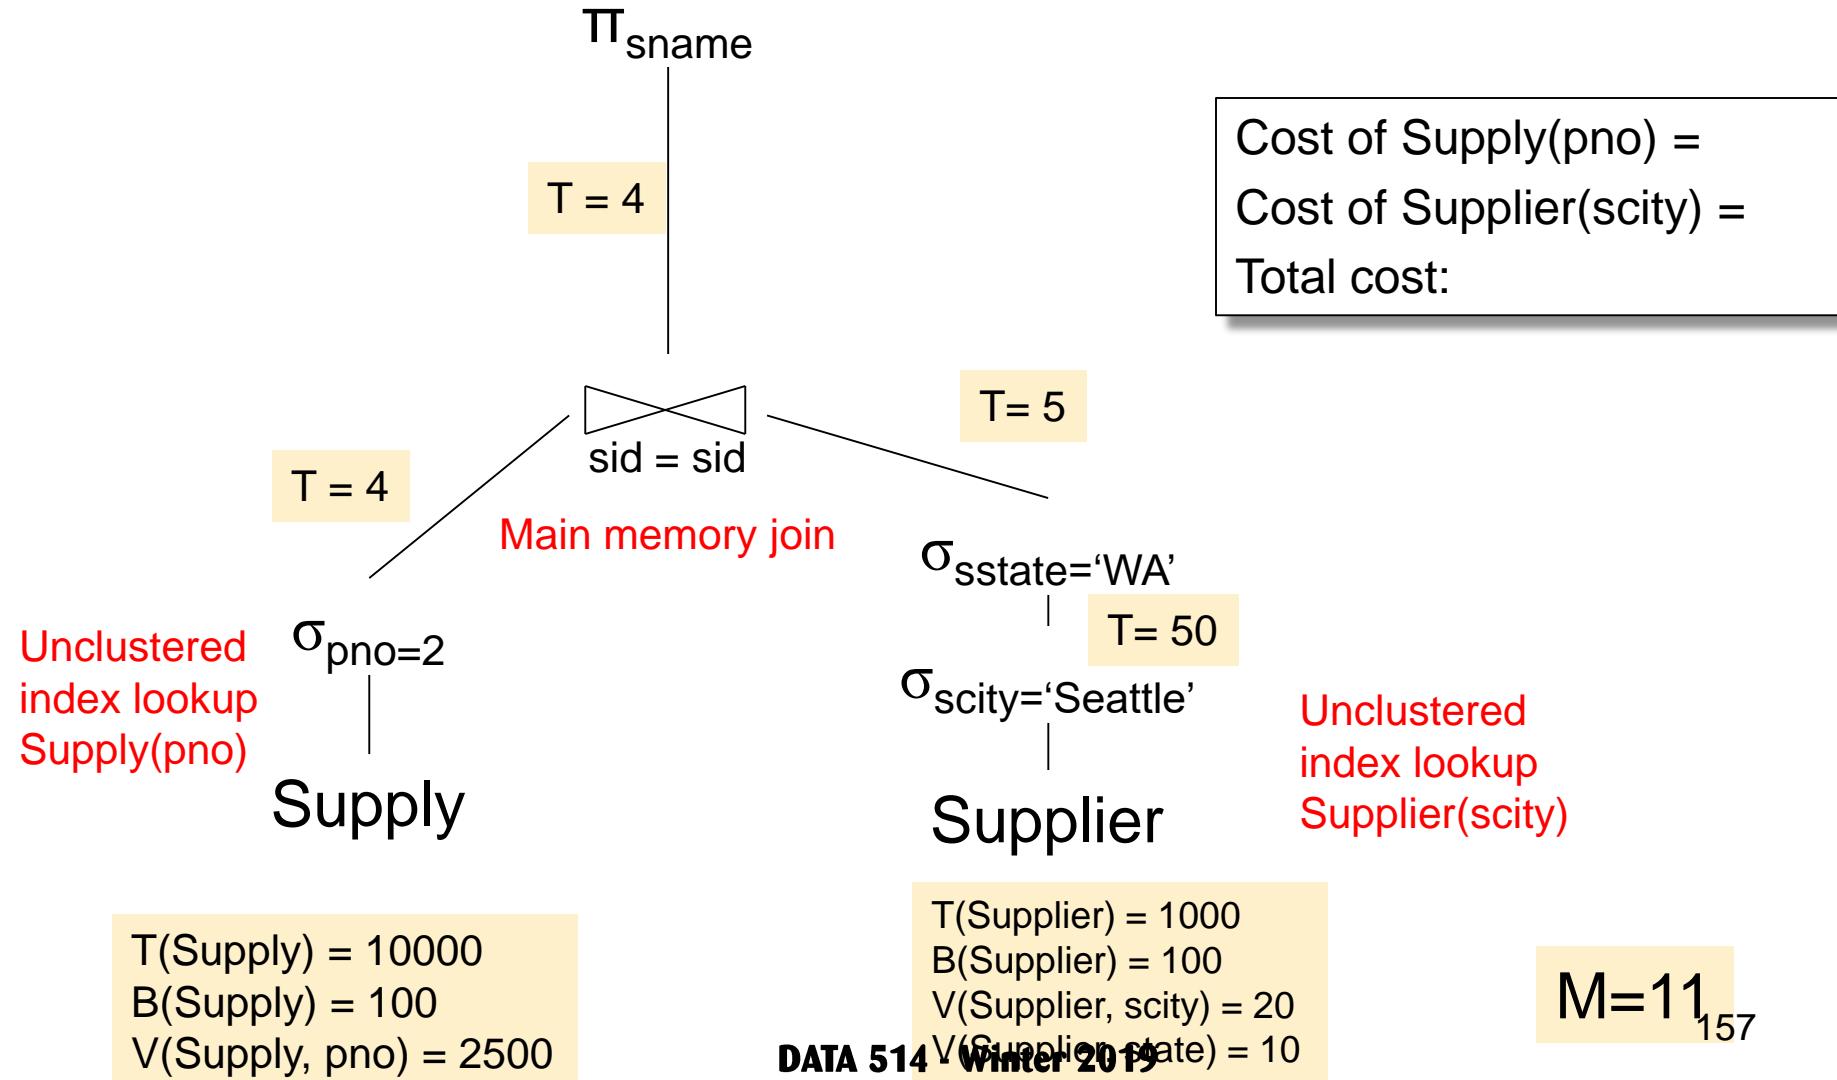
DATA 514 Winter 2019

M=11₁₅₆

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

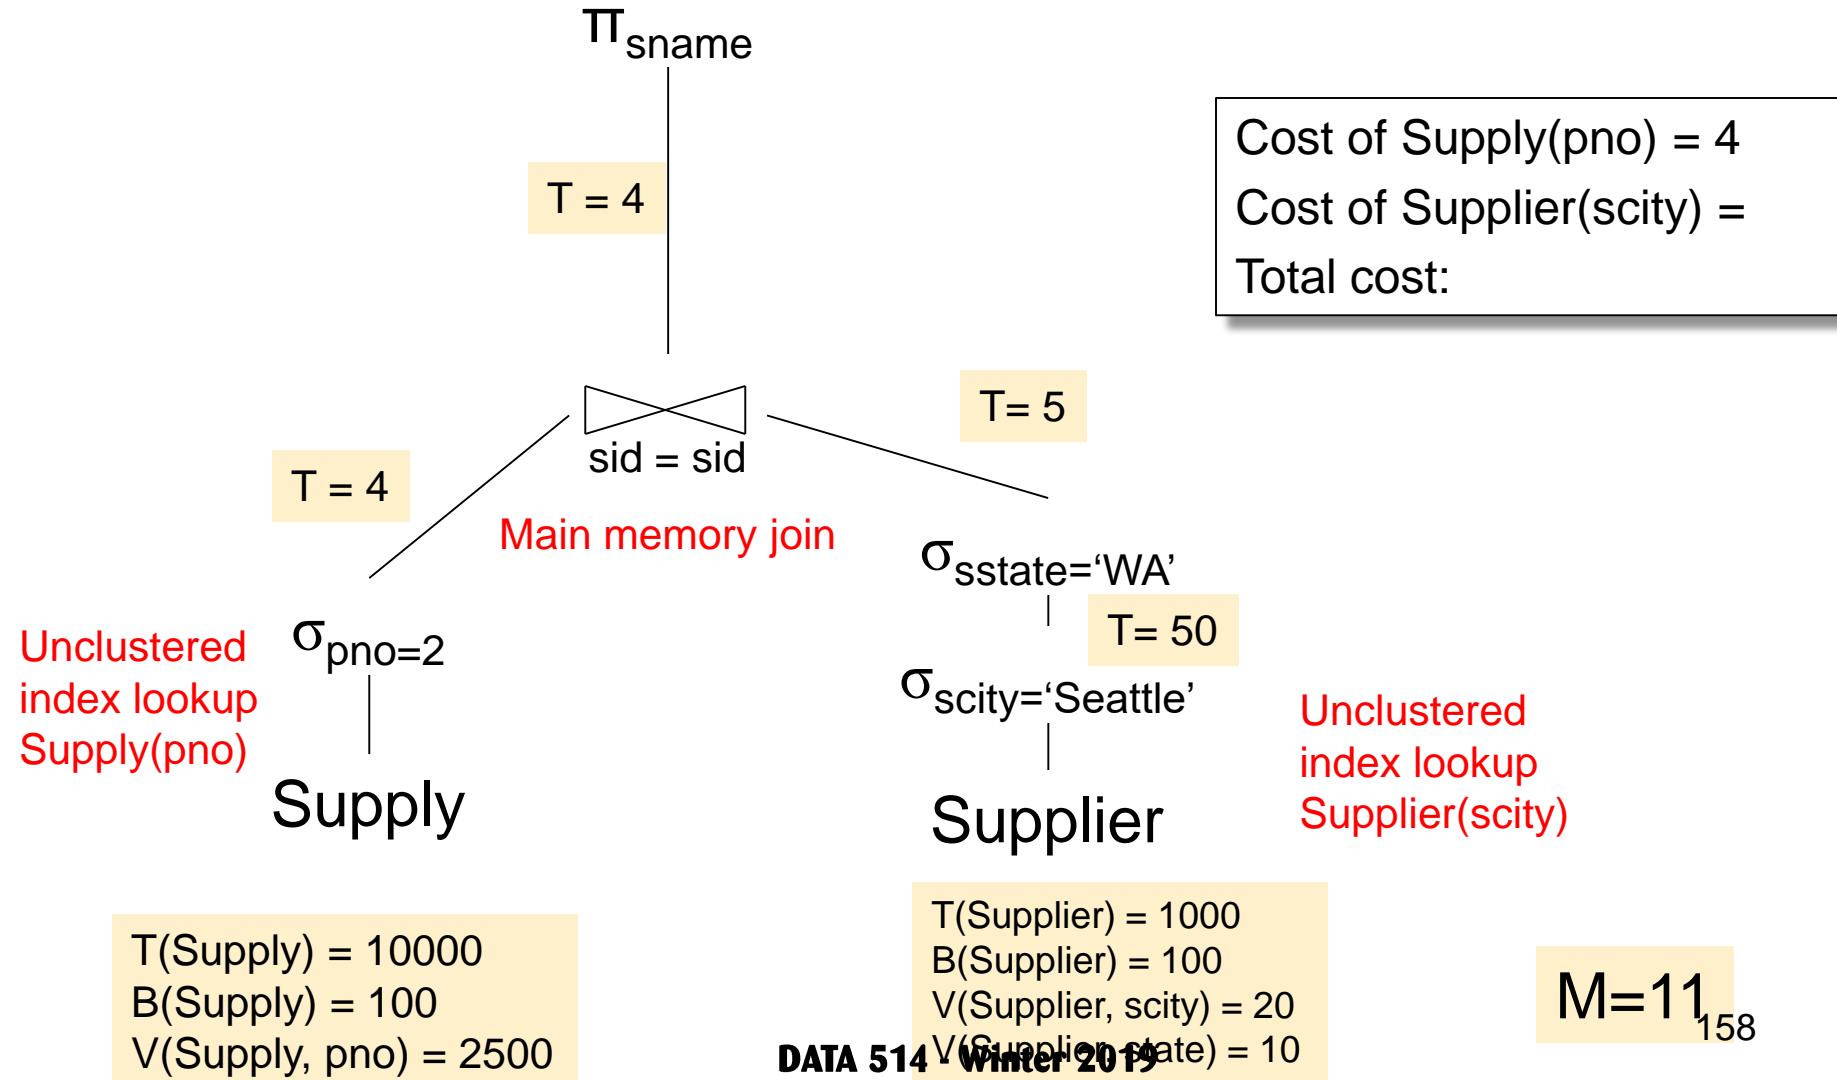
Physical Plan 2



Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

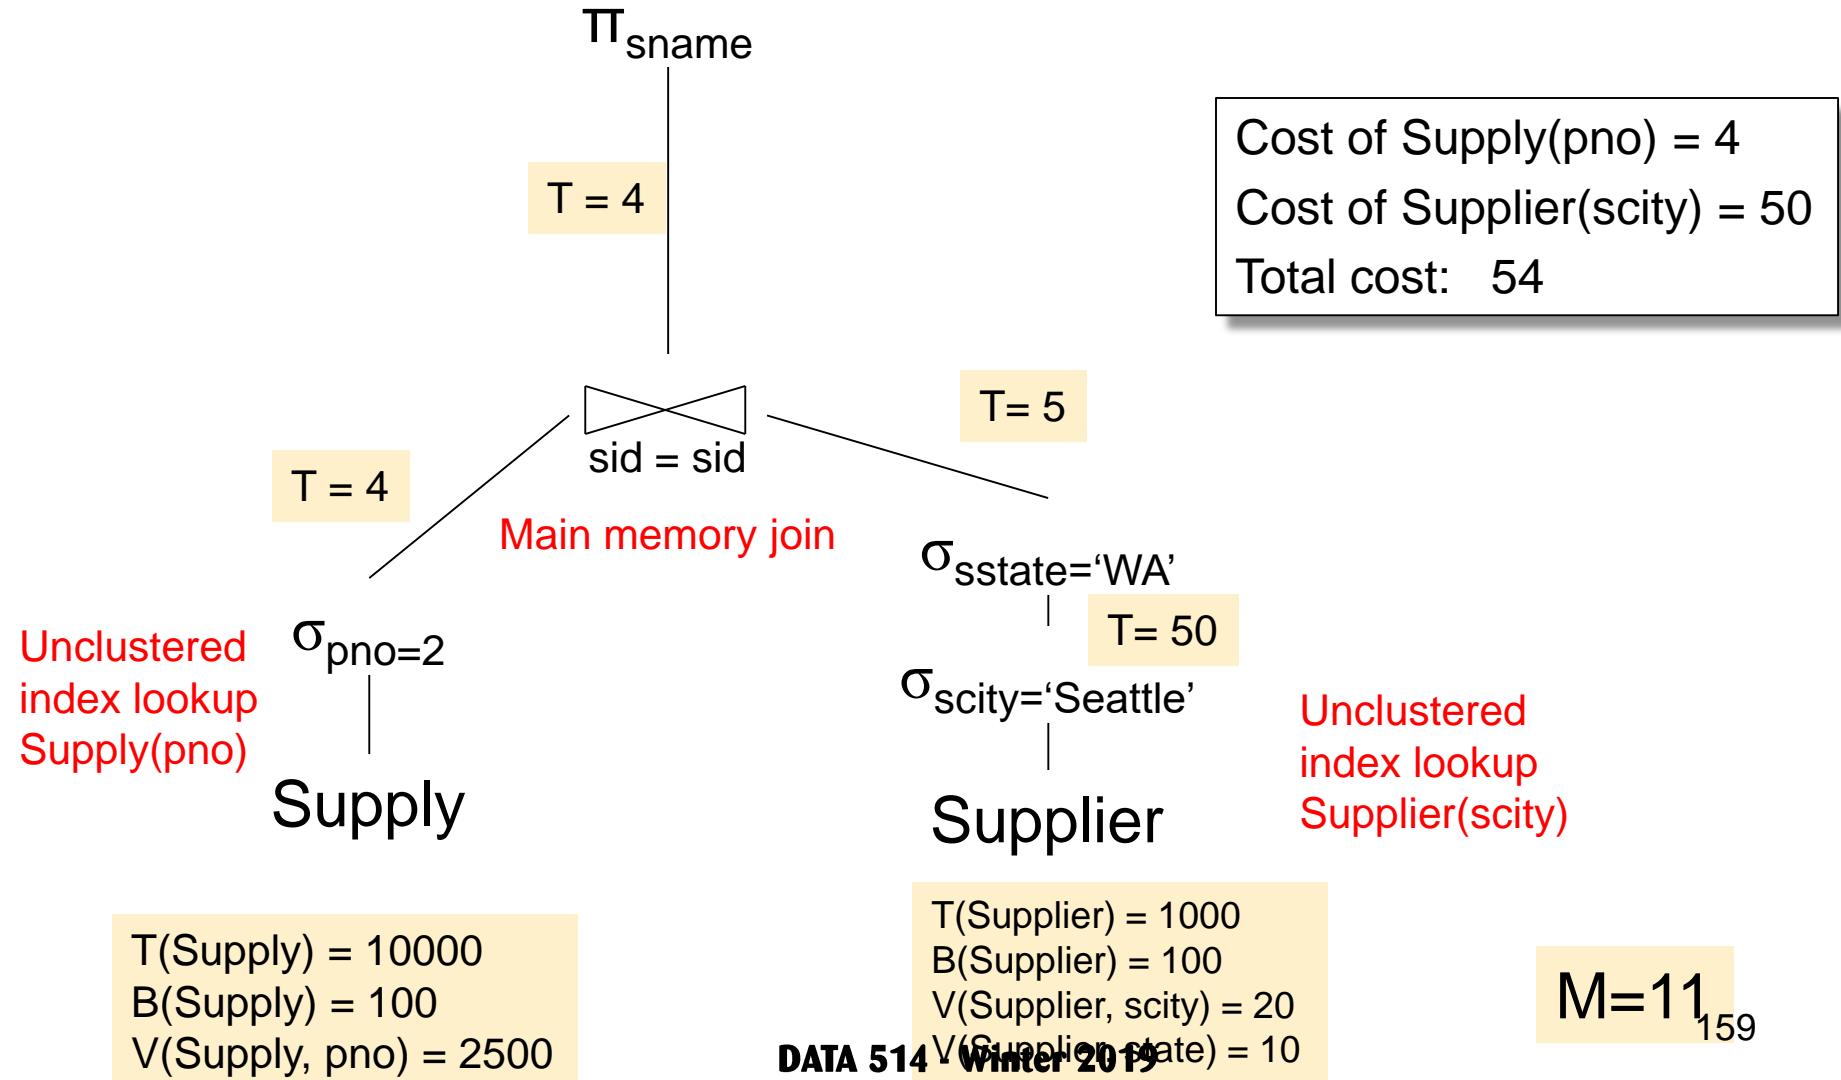
Physical Plan 2



Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

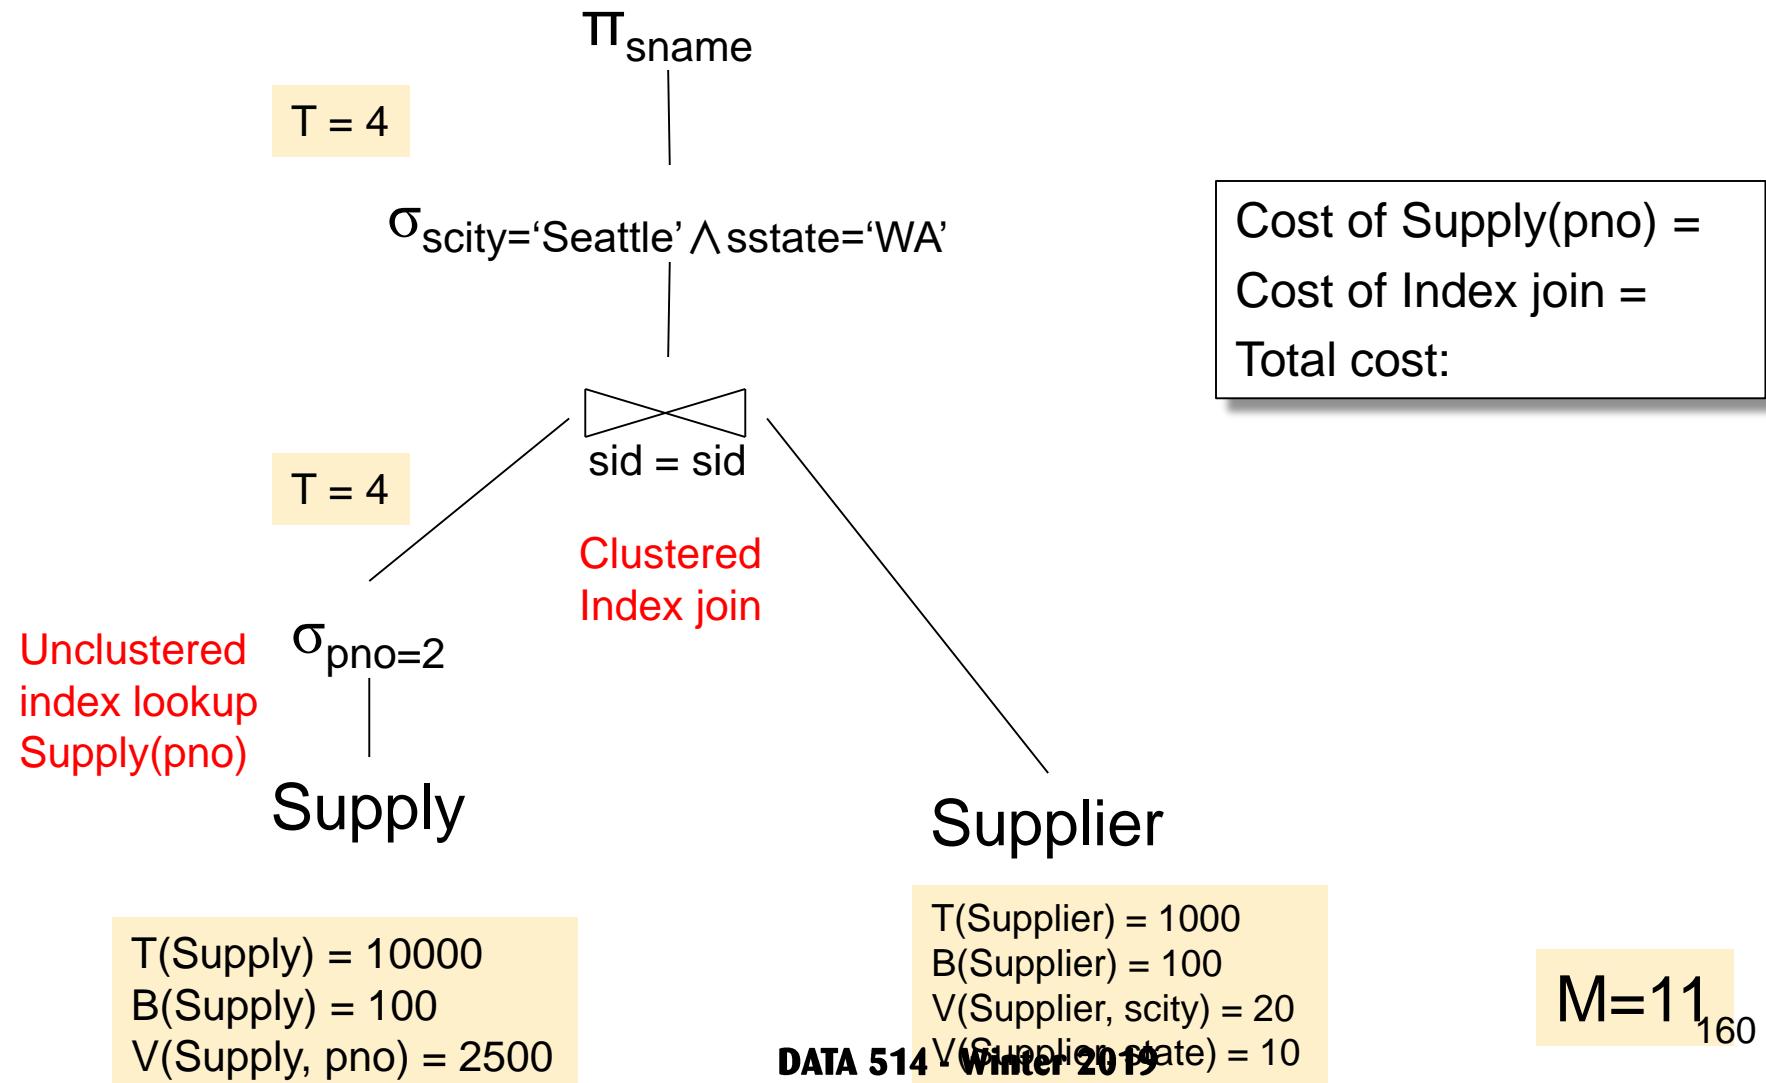
Physical Plan 2



Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

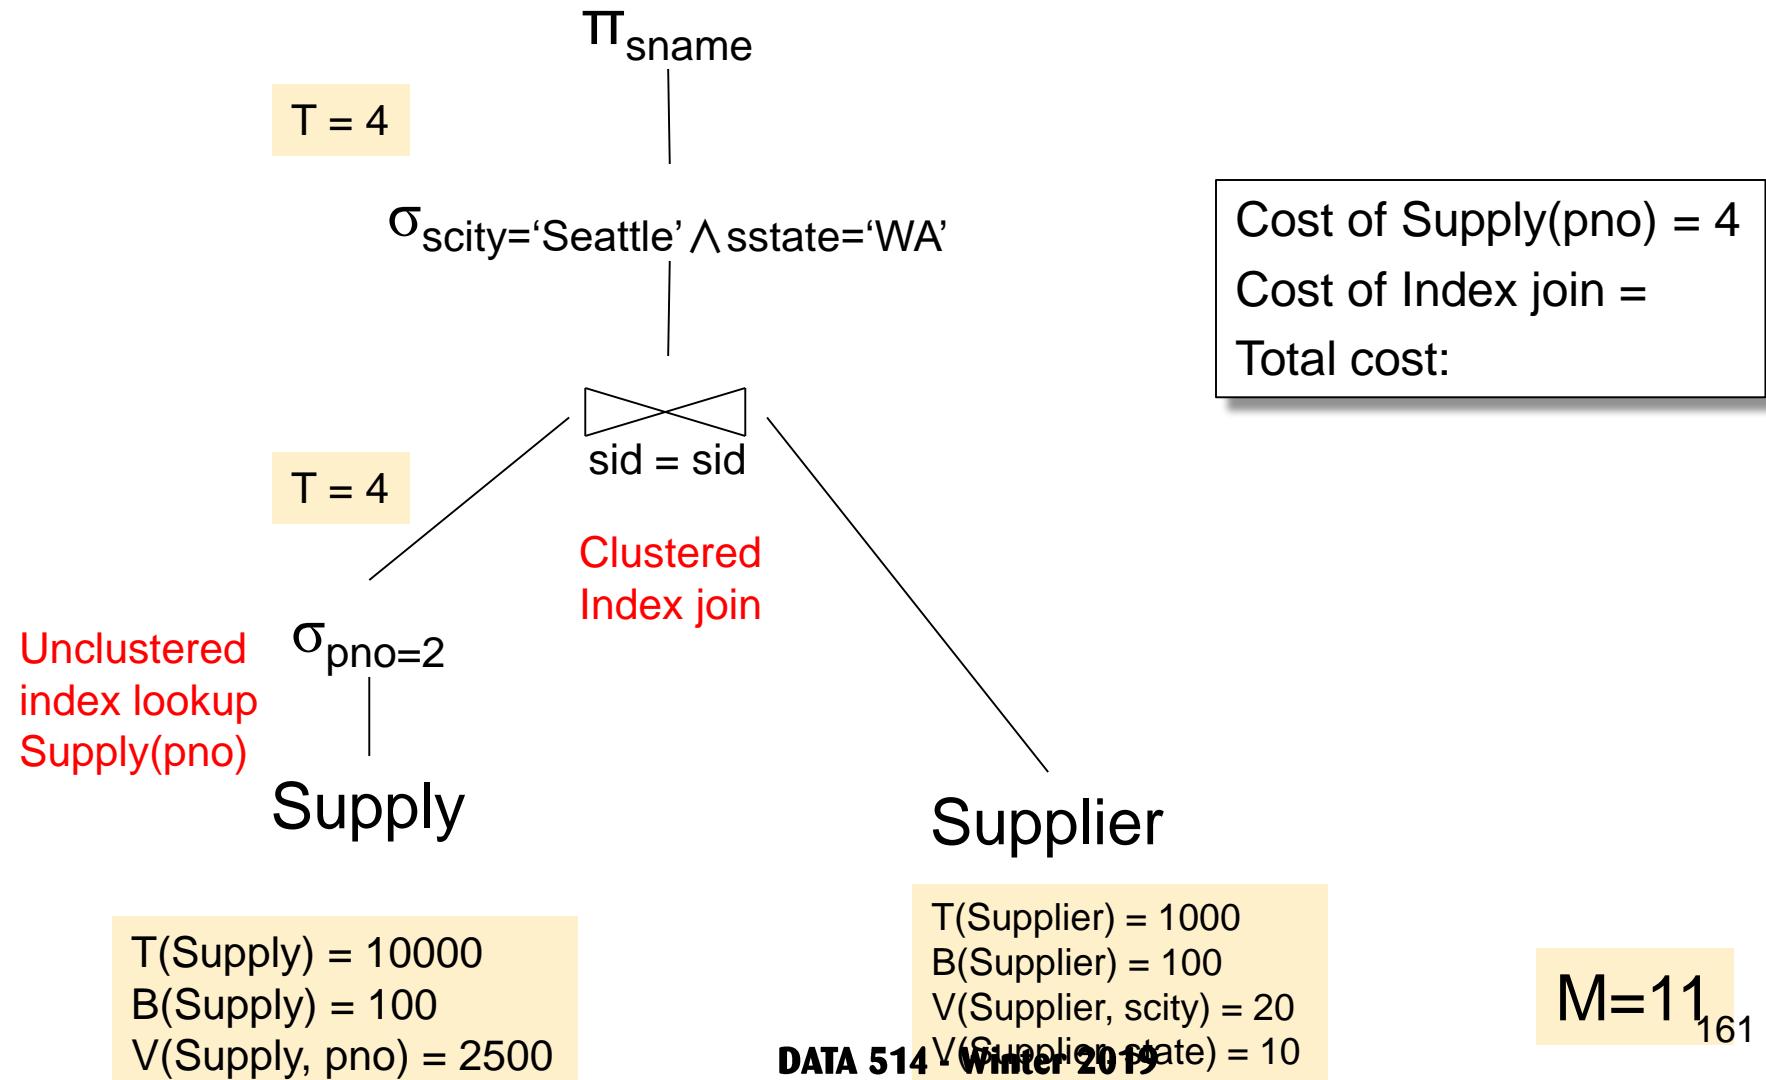
Physical Plan 3



Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

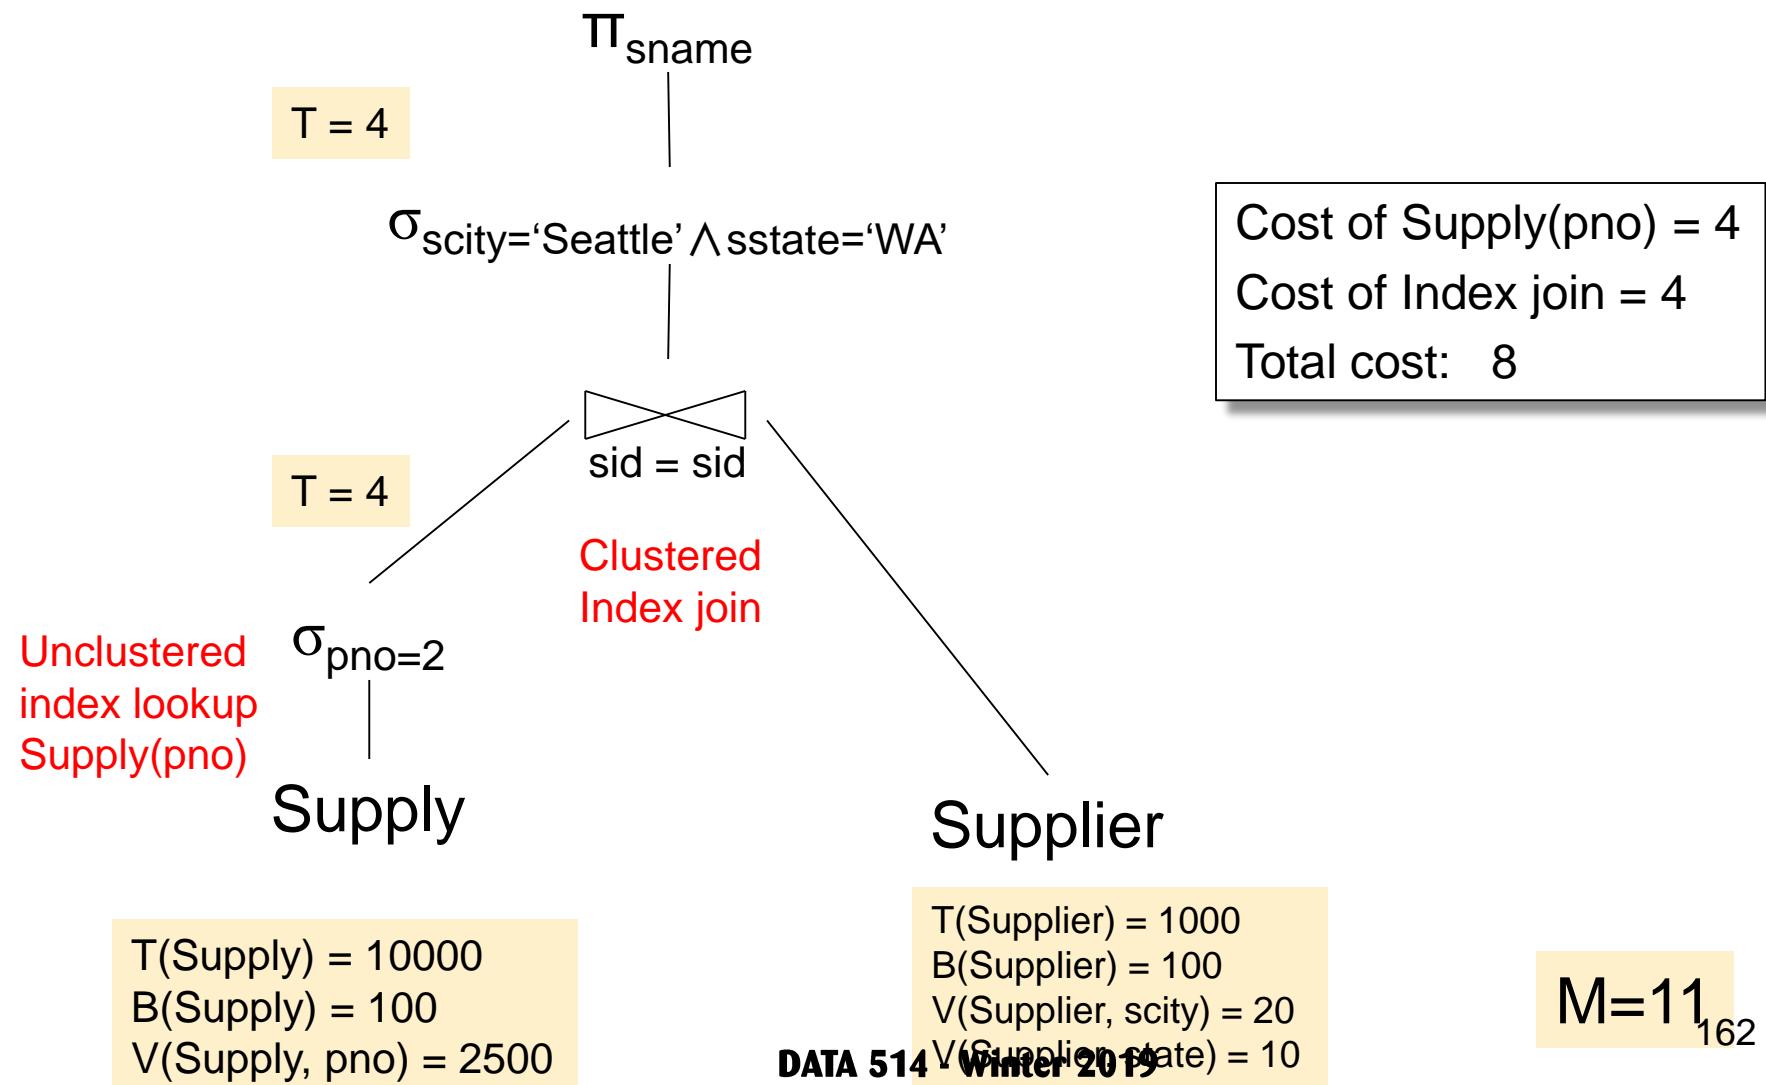
Physical Plan 3



Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Physical Plan 3



Query Optimizer Summary

- **Input: A logical query plan**
- **Output: A good physical query plan**
- **Basic query optimization algorithm**
 - Enumerate alternative plans (logical and physical)
 - Compute estimated cost of each plan
 - Choose plan with lowest cost
- **This is called cost-based optimization**